

Preserving Data Secrecy in Decentralized Inter-organizational Process Mining

Valerio Goretti^a, Davide Basile^a, Luca Barbaro^a, Claudio Di Ciccio^b

^a*Sapienza University of Rome, Via Regina Elena 295, 00161, Rome, Italy*

^b*Utrecht University, Princetonplein 5, 3584 CC, Utrecht, The Netherlands*

Abstract

Inter-organizational business processes involve multiple independent organizations collaborating to achieve mutual interests. Process mining techniques have the potential to allow these organizations to enhance operational efficiency, improve performance, and deepen the understanding of their business based on the recorded process event data. However, inter-organizational process mining faces substantial challenges, including topical secrecy concerns: The involved organizations may not be willing to expose their own data to run mining algorithms jointly with their counterparts or third parties. In this paper, we introduce CONFINE, a novel approach that unlocks process mining on multiple actors' process event data while safeguarding the secrecy and integrity of the original records in an inter-organizational business setting. To ensure that the phases of the presented interaction protocol are secure and that the processed information is hidden from involved and external actors alike, our approach resorts to a decentralized architecture comprised of trusted applications running in Trusted Execution Environments (TEEs). We show the feasibility of our solution by showcasing its application to a healthcare scenario and evaluating our implementation in terms of memory usage and scalability on real-world event logs.

Keywords: Collaborative information

1. Introduction

In today’s business landscape, organizations constantly seek ways to enhance operational efficiency, increase performance, and gain valuable insights to improve their processes. Process mining offers techniques to discover, monitor, and improve business processes by extracting knowledge from chronological records known as *event logs* [1]. Process-aware information systems record events referring to activities and interactions within a business process. The vast majority of process mining contributions consider *intra-organizational* settings, in which business processes are executed inside individual organizations. However, organizations increasingly recognize the value of collaboration and synergy in achieving operational excellence. *Inter-organizational* business processes involve several independent organizations cooperating to achieve a shared objective [2]. Despite the advantages of transparency, performance optimization, and benchmarking that companies can gain, inter-organizational process mining raises challenges that hinder its application. The major issue concerns confidentiality. Companies are reluctant to share private information required to execute process mining algorithms with their partners [3]. Indeed, letting sensitive operational data traverse organizational boundaries introduces concerns about data privacy, security, and compliance with internal regulations [4]. *Trusted Execution Environments* (TEEs) [5] can serve as fundamental enablers to balance the need for insights with the need to protect sensitive information in inter-organizational settings. TEEs offer secure contexts that guarantee code integrity and data confidentiality before, during, and after its utilization.

In this paper, we propose CONFINE, a novel approach and tool aimed at enhanc-

ing collaborative information system architectures with secrecy-preserving process mining capabilities in a decentralized fashion. It resorts to *trusted applications* running in TEEs to preserve the secrecy and integrity of shared data. To pursue this aim, we design a decentralized architecture for a four-staged protocol: (i) The initial exchange of preliminary metadata, (ii) the attestation of the miner entity, (iii) the secure transmission and privacy-preserving merge of encrypted information segments amid multiple parties, (iv) the isolated and verifiable computation of process discovery algorithms on joined data. We evaluate our proof-of-concept implementation against synthetic and real-world-based data with a convergence test followed by experiments to assess the scalability of our approach.

The remainder of this paper is as follows. [Sect. 2](#) provides an overview of related work. In [Sect. 3](#), we introduce a motivating use-case scenario in healthcare. We present the CONFINE approach in [Sect. 5](#). We describe the implementation of our approach in [Sect. 6](#). In [Sect. 7](#), we report on the efficacy and efficiency tests for our solution. Finally, we conclude our work and outline future research directions in [Sect. 8](#).

2. Background and Related Work

2.1. Background

2.2. Related Work

Despite the relative recency of this research branch across process mining and collaborative information systems, scientific literature already includes noticeable contributions to inter-organizational process mining. The work of Müller et al. [6] focuses on data privacy and security within third-party systems that mine data generated from external providers on demand. To safeguard the integrity of data earmarked for mining purposes, their research introduces a conceptual architecture that entails

the execution of process mining algorithms within a cloud service environment, fortified with Trusted Execution Environments. Drawing inspiration from this foundational contribution, our research work seeks to design a decentralized approach characterized by organizational autonomy in the execution of process mining algorithms, devoid of synchronization mechanisms involvement taking place between the involved parties. A notable departure from the framework of Müller et al. lies in the fact that here each participating organization retains the discretion to choose when and how mining operations are conducted. Moreover, we bypass the idea of fixed roles, engineering a peer-to-peer scenario in which organizations can simultaneously be data provisioners or miners. Elkoumy et al. [7, 8] present Shareprom. Like our work, their solution offers a means for independent entities to execute process mining algorithms in inter-organizational settings while safeguarding their proprietary input data from exposure to external parties operating within the same context. Shareprom’s functionality, though, is confined to the execution of operations involving event log abstractions [9] represented as directed acyclic graphs, which the parties employ as intermediate pre-elaboration to be fed into secure multiparty computation (SMPC) [10]. As the authors remark, relying on this specific graph representation imposes constraints that may prove limiting in various process mining scenarios. In contrast, our approach allows for the secure, ciphered transmission of event logs to process mining nodes as a whole. Moreover, SMPC-based solutions require computationally intensive operations and synchronous cooperation among multiple parties, which make these protocols challenging to manage as the number of participants scales up [11]. In our research work, individual computing nodes run the calculations, thus not requiring synchronization with other machines once the input data is loaded.

We are confronted with the imperative task of integrating event logs originat-

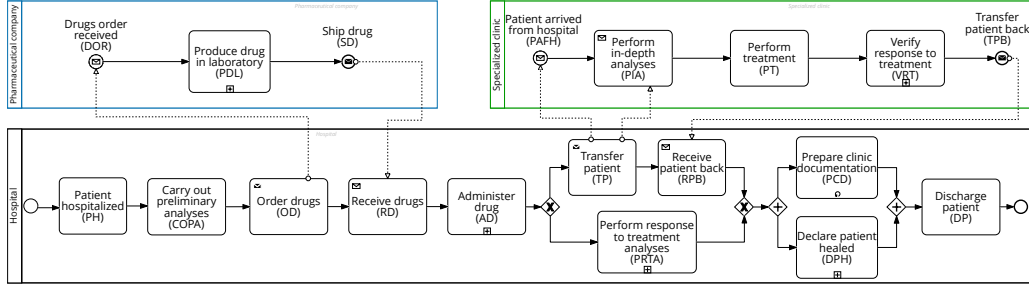


Figure 1: A BPMN collaboration diagram of a simplified healthcare scenario

ing from different data sources and reconstructing consistent traces that describe collaborative process executions. Consequently, we engage in an examination of methodologies delineated within the literature, each of which offers insights into the merging of event logs within inter-organizational settings. The work of Claes et al. [12] holds particular significance for our research efforts. Their seminal study introduces a two-step mechanism operating at the structured data level, contingent upon the configuration and subsequent application of merging rules. Each such rule indicates the relations between attributes of the traces and/or the activities that must hold across distinct traces to be combined. In accordance with their principles, our research incorporates a structured data-level merge based on case references and timestamps as merging attributes. The research by Hernandez et al. [13] posits a methodology functioning at the raw data level. Their approach represents traces and activities as *bag-of-words* vectors, subject to cosine similarity measurements to discern links and relationships between the traces earmarked for combination. An appealing aspect of this approach lies in its capacity to generalize the challenge of merging without necessitating a-priori knowledge of the underlying semantics inherent to the logs under consideration. However, it entails computational overhead in the treatment of data that can interfere with the overall effectiveness of our approach.

Table 1: Events from cases 312 (Alice) and 711 (Bob) recorded by the hospital, the specialized clinic, and the pharmaceutical company

Hospital						Pharmaceutical company			Specialized clinic		
Case	Timestamp	Activity	Case	Timestamp	Activity	Case	Timestamp	Activity	Case	Timestamp	Activity
312	2022-07-14T10:36	PH	312	2022-07-15T22:06	TP	312	2022-07-15T09:06	DOR	312	2022-07-16T00:06	PAFH
312	2022-07-14T16:36	COPA	711	2022-07-16T00:55	PRTA	711	2022-07-15T09:30	DOR	312	2022-07-16T01:06	PIA
711	2022-07-14T17:21	PH	711	2022-07-16T00:55	PCD	312	2022-07-15T11:06	PDL	312	2022-07-16T03:06	PT
312	2022-07-14T17:36	OD	711	2022-07-16T02:55	DPH	711	2022-07-15T11:30	PDL	312	2022-07-16T04:06	VRT
711	2022-07-14T23:21	COPA	711	2022-07-16T04:55	DP	312	2022-07-15T13:06	SD	312	2022-07-16T05:06	TPB
711	2022-07-15T00:21	OD	312	2022-07-16T07:06	RPB	711	2022-07-15T13:30	SD			
711	2022-07-15T18:55	RD	312	2022-07-16T09:06	DPH						
312	2022-07-15T19:06	RD	312	2022-07-16T09:06	PCD						
711	2022-07-15T20:55	AD	312	2022-07-16T11:06	DP						
312	2022-07-15T21:06	AD									

$$T_{312} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, PAFH, PIA, PT, VRT, TPB, RPB, DPH, PCD, DP} \rangle$$

$$T_{711} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, DPH, PCD, DP} \rangle$$

3. Motivating Scenario

For our motivating scenario, we focus on a simplified hospitalization process for the treatment of rare diseases. The process model is depicted as a BPMN diagram in Fig. 1 and involves the cooperation of three parties: a hospital, a pharmaceutical company, and a specialized clinic. For the sake of simplicity, we describe the process through two cases, recorded by the information systems as in Table 2. Alice's journey (case 312) begins when she enters the hospital for the preliminary examinations (patient hospitalized, PH). The hospital then places an order for the drugs (OD) to the pharmaceutical company for treating Alice's specific condition. Afterwards, the pharmaceutical company acknowledges that the drugs order is received (DOR), proceeds to produce the drugs in the laboratory (PDL), and ships the drugs (SD) back to the hospital. Upon receiving the medications, the hospital administers the drug (AD), and conducts an assessment to determine if Alice can be treated internally. If specialized care is required, the hospital transfers the patient (TP) to the specialized clinic. When the patient arrives from the hospital (PAFH), the specialized clinic performs in-depth analyses (PIA) and proceeds with the treatment (PT). Once the specialized clinic had completed the evaluations and verified the response to the treatment (VRT), it transfers the patient back (TPB). The hospital receives the patient

109 back (RPB) and prepares the clinical documentation (PCD). If Alice has successfully
 110 recovered, the hospital declares the patient as healed (DPH). When Alice's treatment
 111 is complete, the hospital discharges the patient (DP). Bob (**case 711**) enters the
 112 hospital a few hours later than Alice. His hospitalization process is similar to Alice's.
 113 However, he does not need specialized care, and his case is only treated by the
 114 hospital. Therefore, the hospital performs the response to treatment analyses (PRTA)
 115 instead of transferring him to the specialized clinic.

116 Both the National Institute of Statistics of the country in which the three organiza-
 117 tions reside and the University that hosts the hospital wish to uncover information on
 118 this inter-organizational process for reporting and auditing purposes [14] via process
 119 analytics. The involved organizations share the urge for such an analysis and wish to
 120 be able to repeat the mining task also in-house. The hospital, the specialized clinic,
 121 and the pharmaceutical company have a partial view of the overall unfolding of the
 122 inter-organizational process as they record the events stemming from the parts of
 123 their pertinence. In Table 1, we show the cases 312 and 711 and the corresponding
 124 traces recorded by the hospital (i.e., T_{312}^H and T_{711}^H), the specialized clinic (i.e., T_{312}^S
 125 and T_{711}^S), and the pharmaceutical company (i.e., T_{312}^C and T_{711}^C). Those traces are
 126 projections of the two combined ones for the whole inter-organizational process:
 127 $T_{312} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, PAFH, PIA, PT, VRT, TPB,}$
 128 $\text{RPB, DPH, PCD, DP} \rangle$ and $T_{711} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP,}$
 129 $\text{DPH, PCD, DP} \rangle$. Results stemming from the analysis of the local cases would not
 130 provide a full picture. Data should be merged. However, to preserve the privacy of
 131 the people involved and safeguard the confidentiality of the information, the involved
 132 parties cannot give open access to their traces to other organizations. The diverging
 133 interests (being able to conduct process mining on data from multiple sources without

134 giving away the local event logs in-clear) motivate our research. In the following,
 135 we describe the design of our solution.

136 4. Notation

137 Given a finite set of events \hat{E} and a total-order relation \leq subset of $\hat{E} \times \hat{E}$, we
 138 identify an event log as the totally ordered set (\hat{E}, \leq) . In the example, ... Let \widehat{CID} be
 139 a finite non-empty set of symbols such that $|\widehat{CID}| \leq |\hat{E}|$. We assume that every event be
 140 associated with a *case identifier* $cid \in \widehat{CID}$ via a total surjective function $\text{id} : \hat{E} \rightarrow \widehat{CID}$
 141 such that the restriction $<_{cid} = \leq \cap \{e \in \hat{E} : \text{id}(e) = cid\}^2$ of total order \leq on all events
 142 mapped to the same cid is strict (i.e., if $e \leq e'$ with $e \neq e'$ and $\text{id}(e) = \text{id}(e')$ then $e' \not\leq e$).
 143 In the example, ... In other words, id acts as an equivalence relation partitioning \hat{E} into
 144 $\{\hat{E}_{cid}\}_{cid \in \widehat{CID}} \subseteq 2^{\hat{E}}$ based on the cid to which the events $e \in \hat{E}_{cid}$ map, and imposing that
 145 events are linearly ordered by the restriction of \leq on every \hat{E}_{cid} . Every pair $(\hat{E}_{cid}, <_{cid})$
 146 thus represents a finite linearly totally ordered set (or *loset* for brevity) with $\hat{E}_{cid} \subseteq \hat{E}$ and
 147 $<_{cid} \subseteq \hat{E}_{cid} \times \hat{E}_{cid} \subseteq \leq \subseteq \hat{E} \times \hat{E}$. Let $(\hat{E}, <)$ be a loset and $(\hat{E}', <')$, $(\hat{E}'', <'')$ two (sub-
 148)losets such that $\hat{E}' \cup \hat{E}'' \subseteq \hat{E}$ and $\hat{E}' \cap \hat{E}'' = \emptyset$, with $<'$ and $<''$ being the restrictions of
 149 $<$ on \hat{E}' and \hat{E}'' , respectively. We define the order-preserving union $\oplus : \hat{E}^3 \times \hat{E}^3 \rightarrow \hat{E}^3$
 150 of losets as follows: $(\hat{E}', <') \oplus (\hat{E}'', <'') = (\hat{E}' \cup \hat{E}'', < \cap (\hat{E}' \cup \hat{E}'')^2)$. We can thus
 151 derive the notion of case C_{cid} given a $cid \in \widehat{CID}$ as a loset of events mapping to
 152 the same cid and ordered by the linear restriction $<$ of \leq over the events in C_{cid} :
 153 $cid = (\hat{E}_{cid}, <)$ where $C_{cid} = \langle e_1, \dots, e_{|C_{cid}|} \rangle$ where $\text{id}(e_i) = cid \in \widehat{CID}$ for every i s.t.
 154 $1 \leq i \leq |C_{cid}|$ and $e_i < e_j$ for every $i \leq j \leq |C_{cid}|$.¹ Notice that the cardinality of \hat{C} and
 155 \widehat{CID} coincide. Events are also the domain of a function $p : \hat{E} \rightarrow \hat{\mathcal{P}}$ mapping events to

Add exam-
ple

Add exam-
ple

Continue
here re-
vising the
definition
of order
preserving
union

¹We employ the angular-bracket notation here for the sake of simplicity, although it is typically used for sequences. Unlike sequences, cases do not allow for the same event to occur more than once.




Figure 2: The CONFINE high-level architecture

156 log provisioners. In the example, . . . We shall denote with $C_{cid}^{\mathcal{P}}$ the loset consisting
 157 of every event $e \in C_{cid}$ such that $p(e) = \mathcal{P}$, with the restriction of the strict total order
 158 of C_{cid} on those events. In the example, . . .

Add exam-
ple

Add exam-
ple

Gotta move this one earlier when we introduce the example with the partitioned event log. (Section 4.2??). Clarify the difference between segmentation (given a segsize, i.e., a segment of a case-part in a sublog) and partitioning (of a log into case-parts of sublogs. Then, prove that the pipeline of partitioning and segmentation has its inverse in the union and merge for soundness.

160 5. Design

161 In this section, we present the high-level architecture of the CONFINE framework.
 162 We consider the main functionalities of each component, avoiding details on the
 163 employed technologies discussed in the next sections.

164 **The CONFINE architecture at large.** Our ar-
 165 chitecture involves different information systems
 166 running on multiple machines. An organiza-
 167 tion can take at least one of the following roles:
 168 **provisioning** if it delivers local event logs to
 169 be collaboratively mined; **mining** if it applies
 170 process mining algorithms using event logs retrieved from provisioners. In Fig. 2,




Figure 3: Sub-components of the Secure Miner

171 we propose the high-level schematization of the CONFINE framework. In our
172 solution, every organization hosts one or more nodes hosting components (the
173 names of which will henceforth be formatted with a teletype font). Depending
174 on the played role, nodes come endowed with a Provisioner or a Secure Miner
175 component, or both. The Provisioner component consists of the following two
176 main sub-components. The Log Recorder registers the events taking place in the
177 organizations' systems. The Log Provider delivers on-demand data to mining
178 players. The hospital and all other parties in our example record Alice and Bob's
179 cases using the Log Recorder. The Log Recorder is queried by the Log Provider
180 for event logs to be made available for mining. The latter controls access to local
181 event logs by authenticating data requests by miners and rejecting those that come
182 from unauthorized parties. In our motivating scenario, the specialized clinic, the
183 pharmaceutical company, and the hospital leverage Log Providers to authenticate
184 the miner party before sending their logs. The Secure Miner component shelters
185 external event logs inside a protected environment to preserve data confidentiality
186 and integrity. Notice that Log Providers accept requests issued solely by Secure
187 Miners. Next, we provide an in-depth focus on the latter.

188 **The Secure Miner.** The primary objective of the Secure Miner is to allow
189 miners to securely execute process mining algorithms using event logs retrieved
190 from provisioners such as the specialized clinic, pharmaceutical company, and the
191 hospital in our example. Secure Miners are isolated components that guarantee
192 data inalterability and confidentiality. In [Fig. 3](#), we show a schematization of
193 the Secure Miner, which consists of four sub-components: (i) Log Requester;
194 (ii) Log Receiver; (iii) Log Manager; (iv) Log Elaborator. The Log Requester
195 and the Log Receiver are the sub-components that we employ during the event log

retrieval. Log Requesters send authenticable data requests to the Log Providers. The Log Receiver collects event logs sent by Log Providers and entrusts them to the Log Manager, securing them from accesses that are external to the Secure Miner. Miners of our motivating scenario, such as the university and the national institute of statistics, employ these three components to retrieve and store Alice and Bob’s data. The Log Elaborator merges the event data locked in the Secure Miner to have a global view of the inter-organizational process comprehensive of activities executed by each involved party. Thereupon, it executes process mining algorithms in a protected environment, inaccessible from the outside computation environment. In our motivating scenario, the Log Elaborator combines the traces associated with the cases of Alice (i.e., T_{312}^H , T_{312}^S , and T_{312}^C) and Bob (i.e., T_{711}^H , T_{711}^S , and T_{711}^C), generates the chronologically sorted traces T_{312} and T_{711} , and feeds them into the mining algorithms (see the bottom-right quadrant of [Table 1](#)).

6. Realization

In this section, we outline the technical aspects concerning the realization of our solution. Therefore, we first present the enabler technologies through which we instantiate the design principles presented in [Sect. 5](#). After that, we discuss the CONFINE interaction protocol. Finally, we show the implementation details.

6.1. Deployment

[Figure 4](#) depicts a UML deployment diagram [15] to illustrate the employed technologies and computation environments. We recall that the Miner and Provisioner nodes are drawn as separated, although organizations can host both. In our motivating scenario, e.g., the hospital can be equipped with machines aimed for both mining and provisioning.

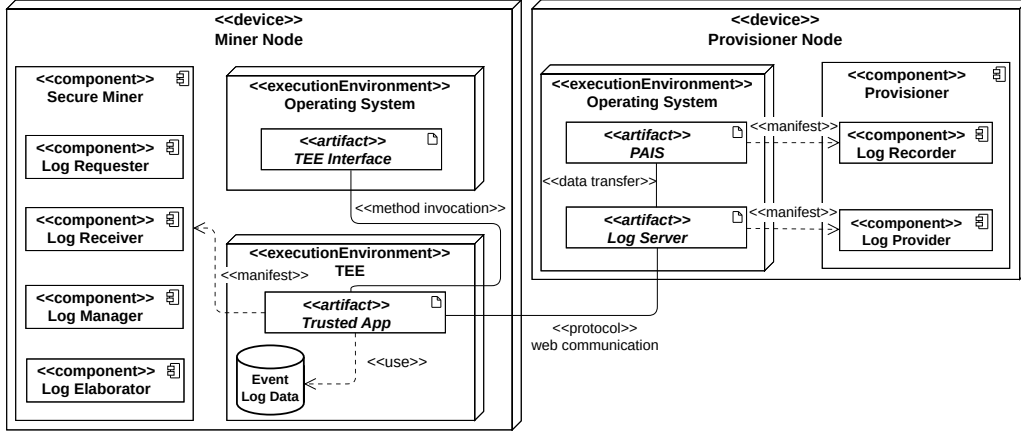


Figure 4: UML deployment diagram of the CONFINE architecture

Provisioner Nodes host the Provisioner’s components, encompassing the Log Recorder and the Log Provider. The Process-Aware Information System (PAIS) manifests the Log Recorder [16]. The PAIS grants access to the Log Server, enabling it to retrieve event log data. The Log Server, on the other hand, embodies the functionalities of the Log Provider, implementing web services aimed at handling remote data requests and providing event log data to miners.

The Miner Node is characterized by two distinct *execution environments*: the Operating System (OS) and the Trusted Execution Environment (TEE) [5]. TEEs establish isolated contexts separate from the OS, safeguarding code and data through hardware-based encryption mechanisms. This technology relies on dedicated sections of a CPU capable of handling encrypted data within a reserved section of the main memory [17]. By enforcing memory access restrictions, TEEs aim to prevent one application from reading or altering the memory space of another, thus enhancing the overall security of the system. This dedicated areas in memory are, however, limited. Once the limits are exceeded, TEEs have to scout around in outer memory areas, thus conceding the opportunity to malicious reader to understand

236 the saved data based on the memory reads and writes. To avoid this risk, TEE
 237 implementations often raise errors that halt the program execution when the memory
 238 demand goes beyond the available space. Therefore, the design of secure systems
 239 that resort to TEEs must take into account that memory consumption must be kept
 240 under control. We leverage the security guarantees provided by TEEs [18] to protect
 241 a Trusted App responsible for fulfilling the functions of the Secure Miner and
 242 its associated sub-components. The TEE ensures the integrity of the Trusted App
 243 code, protecting it against potential malicious manipulations and unauthorized
 244 access by programs running within the Operating System. Additionally, we utilize
 245 the isolated environment of TEEs to securely store event log data (e.g., Alice and
 246 Bob’s cases). The TEE retains a private key in the externally inaccessible section of
 247 memory, paired with a public key in a Rivest-Shamir-Adleman (RSA) [19] scheme
 248 for attestation (only the owner of the private key can sign messages that are verifiable
 249 via the public key) and secure message encryption (only the owner of the private key
 250 can decode messages that are encrypted with the corresponding public key). In our
 251 solution, access to data located in the TEE is restricted solely to the Trusted App.
 252 Users interact with the Trusted App through the Trusted App Interface, which
 253 serves as the exclusive communication channel. The Trusted App offers secure
 254 methods, invoked by the Trusted App Interface, for safely receiving information
 255 from the Operating System and outsourcing the results of computations.

256 6.2. *The CONFINE protocol*

257 We orchestrate the interaction of the components in CONFINE via a protocol.

258 Our protocol involves two main actors: a Secure Miner (henceforth, \mathcal{M}) and
 259 one to many Log Provisioners ($\mathcal{P}_1, \dots, \mathcal{P}_n \in \hat{\mathcal{P}}$). Their behaviour is documented in
 260 [Alg. 1](#) and [Alg. 2](#)




Figure 5: Unfolding example for the initialization, remote attestation phases of the CONFINE protocol

261 We separate it in four subsequent stages, namely (i) *initialization*, (ii) *remote*
 262 *attestation*, (iii) *data transmission*, and (iv) *computation*. These stages are depicted
 263 in Figs. 5(a), 5(b), 6(a) and 6(b), respectively. They are mainly enacted by a Miner
 264 Node and n Provisioner Nodes. We assume their communication channel is
 265 reliable [20] and secure [21]. In the following, we describe each phase in detail.

266 **Initialization.** The objective of the initialization stage is to inform the miner about the
 267 distribution of cases related to a business process among the Provisioner Nodes.
 268 At the onset of this stage, the Log Requester within the Trusted App issues n
 269 requests, one per Log Server component, to retrieve the list of case references they
 270 record (step 1 in Fig. 5(a)). Following sender authentication (2), each Log Server
 271 retrieves the local event log from the PAIS (3, 4) and subsequently responds to the Log
 272 Requester by providing a list of its associated case references (5). After collecting
 273 these n responses, the Log Requester delineates the distribution of cases. In the

context of our motivating scenario, by the conclusion of the initialization, the miner gains knowledge that the case associated with Bob, synthesized in the traces T_{711}^H and T_{711}^C , is exclusively retained by the hospital and the specialized clinic. In contrast, the traces of Alice’s case, denoted as T_{312}^H , T_{312}^C , and T_{312}^S , are scattered across all three organizations.

Remote attestation. The remote attestation serves the purpose of establishing trust between miners and provisioners in the context of fulfilling data requests. This phase adheres to the overarching principles outlined in the RATS RFC standard [22] serving as the foundation for several TEE attestation schemes (e.g., Intel EPID,² and AMD SEV-SNP³). Remote attestation has a dual objective: (i) to furnish provisioners with compelling evidence that the data request for an event log originates from a Trusted App running within a TEE; (ii) to confirm the specific nature of the Trusted App as an authentic Secure Miner software entity. This phase is triggered when the Log Requester sends a new case request to the Log Server, specifying: (i) the segment size (henceforth, *seg_size*), and (ii) the set of the requested case references. Both parameters will be used in the subsequent *data transmission* phase. Each of the n Log Servers commences the verification process by requesting the necessary information from the Log Receiver to conduct the attestation (2). Subsequently, the Log Receiver generates the attestation report containing the so-called *measurement* of the Trusted App, which is defined as the hash value of the combination of its source code and data. Once this report is signed using the attestation private key associated with the TEE’s hardware of the Miner Node, it is transmitted by the Log

²sgx101.gitbook.io/sgx101/sgx-bootstrap/attestation. Accessed: 15/01/2024.

³amd.com/en/processors/amd-secure-encrypted-virtualization. Accessed: 15/01/2024.

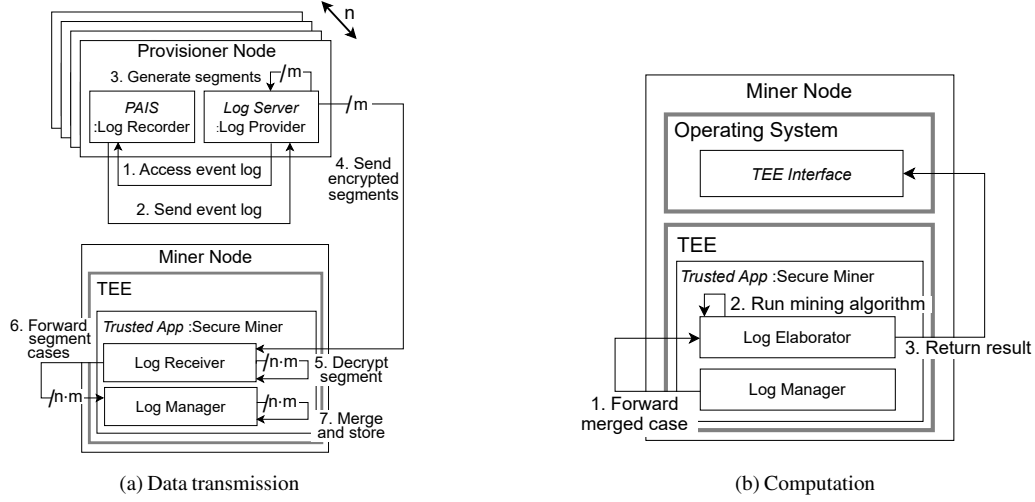


Figure 6: Unfolding example for the data transmission and computation phases of the CONFINE protocol

Receiver to the Log Servers alongside the attestation public key of the Miner Node (3). The Log Servers authenticate the miner using the public key and decrypt the report (4). In this last step, the Log Servers undertake a comparison procedure in which they juxtapose the measurement found within the decrypted report against a predefined reference value associated with the source code of the Secure Miner. If the decrypted measurement matches the predefined value, the Miner Node gains trust from the provisioner.

Data transmission. Once the trusted nature of the Trusted App is verified, the Log Servers proceed with the transmission of their cases. To accomplish this, each Log Server retrieves the event log from the PAIS (steps 1 and 2, in Fig. 6(a)), and filters it according to the case reference set specified by the miner. Given the constrained workload capacity of the TEE, it is imperative for Log Servers to partition the filtered event log into distinct segments. Consequently, each Log Server generates m log segments comprising a variable count of entire cases (3). The cumulative size of

310 these segments is governed by the threshold parameter specified by the miner in the
 311 initial request (step 1 of the remote attestation phase, Fig. 5(b)). As an illustrative
 312 example from our motivating scenario, the Log Server of the hospital may structure
 313 the segmentation such that T_{312}^H and T_{711}^H reside within the same segment, whereas
 314 the specialized clinic might have T_{312}^S and T_{711}^S in separate segments. Subsequently,
 315 the n Log Servers transmit their m encrypted segments to the Log Receiver of
 316 the Trusted App (4). The Log Receiver, in turn, collects the $n \times m$ responses in a
 317 queue, processing them one at a time. After decrypting the processed segment (5),
 318 the Log Receiver forwards the cases contained in the segment to the Log Manager
 319 (6). To reconstruct the process instance, cases belonging to the same process instance
 320 must be merged by the Log Manager resulting in a single trace (e.g., T_{312} for Alice
 321 case) comprehensive of all the events in the partial traces (e.g., T_{312}^H , T_{312}^S and T_{312}^C
 322 for Alice case). During this operation, the Log Manager applies a specific *merging*
 323 *schema* (i.e., a rule specifying the attributes that link two cases during the merge)
 324 as stated in [12]. In our illustrative scenario, the merging schema to combine the
 325 cases of Alice is contingent upon the linkage established through their case identifier
 326 (i.e., 312). We underline that our proposed solution facilitates the incorporation
 327 of diverse merging schemas encompassing distinct trace attributes. The outcomes
 328 arising from merging the cases within the processed segment are securely stored by
 329 the Log Manager in the TEE.

330 **Computation.** The Trusted App requires all the provisioners to have delivered cases
 331 referring to the same process instances. For example, when the hospital and the other
 332 organizations have all delivered their information concerning case 312 to the Trusted
 333 App, the process instance associated with Alice becomes eligible for computation.
 334 Upon meeting this condition, the Log Manager forwards the case earmarked for

Algorithm 1: Secure Miner's behavior in CONFINE.

Input: $\hat{\mathcal{P}} = \{\mathcal{P}_1, \dots, \mathcal{P}_n\}$, the (references to) n log provisioners;
 seg_size , the maximum size of the log segment to be transmitted by the log provisioners.

Data: $CIDMap: \widehat{CID} \rightarrow 2^{\hat{\mathcal{P}}}$, a map from case references $cid \in \widehat{CID}$ to the set of log provisioners in $\hat{\mathcal{P}}$;
 $PMap: \hat{\mathcal{P}} \rightarrow 2^{\widehat{CID}}$, a map from log provisioners $\mathcal{P} \in \hat{\mathcal{P}}$ to the set of references to their cases in \widehat{CID} ;
 $Cases: \widehat{CID} \rightarrow \hat{\mathcal{C}}$, a map from case references $cid \in \widehat{CID}$ to a set of cases in $\hat{\mathcal{C}}$.

Implements: `SecureMiner`, instance \mathcal{M} .

Uses: `AuthenticatedPerfectPointToPointLink`, instance aI .

```

1  upon event  $\langle \mathcal{M}, \text{INIT} | \hat{\mathcal{P}}, seg\_size \rangle$  do           // The Log Requester of  $\mathcal{M}$  starts the CONFINE protocol – initialization phase in Fig. 5(a)
2    foreach  $\mathcal{P} \in \hat{\mathcal{P}}$  do                               // For every Provisioner  $\mathcal{P}$ 
3      trigger  $\langle aI, \text{SEND} | \mathcal{P}, [\text{CASESREFREQ}] \rangle$       // Request  $\mathcal{P}$ 's case references (see Alg. 2, line 1)
4    upon  $|\hat{\mathcal{P}}| = |\text{dom}(PMap)|$  do                       // Once all Provisioners have answered with their case references
5      foreach  $\mathcal{P} \in \hat{\mathcal{P}}$  do trigger  $\langle aI, \text{SEND} | \mathcal{P}, [\text{CASESREQ}, seg\_size, PMap[\mathcal{P}]] \rangle$  // Request their cases via  $aI$  (see Alg. 2, line 4)

6  upon event  $\langle aI, \text{DELIVER} | \mathcal{P}, [\text{CASESREFRES}, CIDs] \rangle$  such that  $\mathcal{P} \in \hat{\mathcal{P}}$  do //  $\mathcal{M}$ 's Log Requester gets  $\mathcal{P}$ 's case references via  $aI$  (Alg. 2, line 3)
7    foreach  $cid \in CIDs$  do                             // For every case reference  $cid$  received in  $CIDs$ 
8       $CIDMap[cid] \leftarrow CIDMap[cid] \cup \{\mathcal{P}\}$       // Add  $\mathcal{P}$  to the set of provisioners for case  $cid$  in  $CIDMap$ 
9       $PMap[\mathcal{P}] \leftarrow PMap[\mathcal{P}] \cup CIDs$            // Register the references of the cases provided by  $\mathcal{P}$  in  $PMap$ 

10 upon event  $\langle aI, \text{DELIVER} | \mathcal{P}, [\text{CASESRES}, S] \rangle$  such that  $\mathcal{P} \in \hat{\mathcal{P}}$  do //  $\mathcal{M}$ 's Log Receiver gets a segment from  $\mathcal{P}$  via  $aI$  (Alg. 2, line 8)
11   foreach  $C_{cid}^{\mathcal{P}} \in S$  do                          // For every  $C_{cid}^{\mathcal{P}}$  in the delivered segment  $S$ , each associated with a  $cid$ – data transmission phase in Fig. 6(a)
12     if  $cid \in PMap[\mathcal{P}]$  then                          // If  $\mathcal{P}$  has declared the ownership of  $cid$  (see line 6)
13        $PMap[\mathcal{P}] \leftarrow PMap[\mathcal{P}] \setminus \{cid\}$       // Remove  $cid$  from the set of case references to be provided by  $\mathcal{P}$ 
14        $CIDMap[cid] \leftarrow CIDMap[cid] \setminus \{\mathcal{P}\}$  // Remove  $\mathcal{P}$  from the set of  $cid$  provisioners
15        $\mathcal{M}.\text{LogManager.mergeAndStore}(Cases, C_{cid}^{\mathcal{P}})$  // Update the case via  $\oplus$  and store the result in  $Cases$ 

16 upon  $CIDMap[cid] = \emptyset$  for some  $cid \in \text{dom}(CIDMap)$  do // When all the pieces of some  $cid$  have arrived to  $\mathcal{M}$ 's Log Manager
17    $\text{dom}(CIDMap) \leftarrow \text{dom}(CIDMap) \setminus \{cid\}$  // Remove  $cid$  from the domain of cases which still needs to be processed
18   yield  $Cases[cid]$  to  $\mathcal{M}.\text{LogElaborator}$  // Forward the case  $cid$  to the Log Elaborator of  $\mathcal{M}$  for mining – computation phase in Fig. 6(b)

```

335 computation to the Log Elaborator (step 1 in Fig. 6(b)). Subsequently, the Log
336 Elaborator proceeds to input the merged case into the process mining algorithm
337 (2). Ultimately, the outcome of the computation is relayed by the Log Elaborator
338 from the TEE to the TEE Interface running atop the Operating System of the
339 Miner Node (3). The CONFINE protocol does not impose restrictions on the
340 post-computational handling of results. In our motivating scenario, the University
341 and the National Institute of Statistics, serving as miners, disseminate the outcomes of
342 computations, generating analyses that benefit the provisioners (though the original
343 data are never revealed in clear). Furthermore, our protocol enables the potential for
344 provisioners to have their proprietary Secure Miner, allowing them autonomous
345 control over the computed results.

Algorithm 2: Provisioner’s behavior in CONFINE.

Input: $\widehat{\mathcal{M}} = \{\mathcal{M}_1, \dots, \mathcal{M}_s\}$, the (references to) s miners.
Implements: Provisioner, instance \mathcal{P} .
Uses: *AuthenticatedPerfectPointToPointLink*, instance aL .

```

1 upon event  $\langle aL, \text{DELIVER} | \mathcal{M}, [\text{CASESREFSREQ}] \rangle$  such that  $\mathcal{M} \in \widehat{\mathcal{M}}$  do //  $\mathcal{P}$  receives the request for case references from  $\mathcal{M}$  (see Alg. 1, line 3)
2    $CIDs \leftarrow \mathcal{P}.\text{LogRecorder}.\text{accessCaseReferences}()$  // Access the case references via Log Recorder
3   trigger  $\langle aL, \text{SEND} | \mathcal{M}, [\text{CASESREFRES}, CIDs] \rangle$  // send the case references to  $\mathcal{M}$  (see Alg. 1, line 6)

4 upon event  $\langle aL, \text{DELIVER} | \mathcal{M}, [\text{CASESREQ}, \text{seg\_size}, CIDs] \rangle$  such that  $\mathcal{M} \in \widehat{\mathcal{M}}$  do //  $\mathcal{P}$  gets the case request from  $\mathcal{M}$  (see Alg. 1, line 5)
5   if  $\mathcal{M}.\text{LogReceiver}.\text{getAttestationReport}(\mathcal{P})$  is valid then // Get and verify the attestation report of  $\mathcal{M}$  – remote attestation in Fig. 5(b)
6      $\{S_1, \dots, S_m\} \leftarrow \mathcal{P}.\text{LogProvider}.\text{segmentEventLog}(\mathcal{P}.\text{LogRecorder}.\text{accessEventLog}(CIDs), \text{seg\_size})$  // Segment the event log
7     foreach  $i \in \{1, \dots, m\}$  do // For every split segment  $S_i$ 
8       trigger  $\langle aL, \text{SEND} | \mathcal{M}, [\text{CASESRES}, S_i] \rangle$  // send the segment  $S_i$  to  $\mathcal{M}$  (see Alg. 1, line 10) – data transmission phase in Fig. 6(a)
  
```

Table 2: Event logs used for our experiments

Name	Type	Activities	Cases	Max events	Min events	Avg. events	Organization \mapsto Activities
Motivating scenario	Synthetic	19	1000	18	9	14	$\mathcal{O}^P \mapsto 3, \mathcal{O}^C \mapsto 5, \mathcal{O}^H \mapsto 14$
Sepsis [25]	Real	16	1050	185	3	15	$\mathcal{O}^1 \mapsto 1, \mathcal{O}^2 \mapsto 1, \mathcal{O}^3 \mapsto 14$
BPIC2013 [26]	Real	7	1487	123	1	9	$\mathcal{O}^1 \mapsto 6, \mathcal{O}^2 \mapsto 7, \mathcal{O}^3 \mapsto 6$

6.3. Implementation

We implemented the Secure Miner component as an Intel SGX⁴ trusted application, encoded in Go through the EGo framework.⁵ We resort to a TLS [23] communication channel between miners and provisioners over the HTTP web protocol to secure the information exchange. To demonstrate the effectiveness of our framework, we re-implemented and integrated the *HeuristicsMiner* discovery algorithm [24] within the Trusted Application. Our implementation of CONFINE, including the *HeuristicsMiner* in Go, is openly accessible at the following URL: github.com/Process-in-Chains/CONFINE/.

⁴sgx101.gitbook.io/sgx101/. Accessed: 15/01/2024.

⁵docs.edgeless.systems/ego. Accessed: 15/01/2024.




Figure 7: *HeuristicsMiner* output in CONFINE

7. Evaluation

In this section, we evaluate our approach through the testing of our tool implementation. We begin with a convergence analysis to demonstrate the correctness of the collaborative data exchange process. Subsequently, we gauge the memory usage with synthetic and real-life event logs, to observe the trend during the enactment of our protocol and assess scalability. We recall that we focus on memory utilization since the availability of space in the dedicated areas is limited as we discussed in Sect. 6.1. We discuss our experimental results in the following. For the sake of reproducibility, we make available all the testbeds and results in our public code repository (linked above).

Output convergence. To experimentally validate the correctness of our approach in the transmission and computation phases (see Sect. 6), we run a *convergence* test. To this end, we created a synthetic event log consisting of 1000 cases of 14 events on average (see Table 2) by simulating the inter-organizational process of our motivating scenario (see Fig. 1)⁶ and we partitioned it in three sub-logs (one per involved organization), an excerpt of which is listed in Table 1. We run the stand-alone *HeuristicsMiner* on the former, and processed the latter through our

⁶We generated the event log through BIMP (<https://bimp.cs.ut.ee/>). We filtered the generated log by keeping the sole events that report on the completion of activities, and removing the start and end events of the pharmaceutical company and specialized clinic’s sub-processes.




Figure 8: Memory usage test results

CONFINE toolchain. As expected, the results converge and are depicted in Fig. 7 in the form of a workflow net [27]. For clarity, we have colored activities recorded by the organizations following the scheme of Table 2 (black for the hospital, blue for the pharmaceutical company, and green for the specialized clinic).

Memory usage. Figures 8(a) and 8(b) display plots corresponding to the runtime memory utilization of our CONFINE implementation (in MegaBytes). Differently from Fig. 8(b), Fig. 8(a) excludes the computation stage by leaving the *HeuristicsMiner* inactive so as to isolate the execution from the mining-specific operations. The dashed lines mark the starting points for the remote attestation, the data transmission and the computation stages. We held the *seg_size* constant at 2000 KiloBytes. We observe that the data transmission stage reaches the highest peak of memory utilization, which is then partially freed by the subsequent computation

stage, steadily occupying memory space at a lower level. To verify whether this phenomenon is due to the synthetic nature of our simulation-based event log, we also gauge the runtime memory usage of two public real-world event logs too (Sepsis [25] and BPIC 2013 [26]). The characteristics of the event logs are summarized in Table 2. Since those are *intra-organizational* event logs, we split the contents to mimic an *inter-organizational* context. In particular, we separated the Sepsis event log based on the distinction between normal-care and intensive-care paths, as if they were conducted by two distinct organizations. Similarly, we processed the BPIC 2013 event log to sort it out into the three departments of the Volvo IT incident management system. Figure 8(c) depicts the results. We observe that the processing of the BPIC 2013 event log demands more memory, particularly during the initial stages, probably owing to its larger size. Conversely, the Sepsis event log turns out to entail the least expensive run. To verify whether these trends are affected by the dimension of the exchanged data segments, we conducted an additional test to examine the trend of memory usage as the *seg_size* varies with all the aforementioned event logs. Notably, the polylines displayed in Fig. 8(d) indicate a linear increment of memory occupation until a breakpoint is reached. After that, the memory in use is steady. These points, marked by vertical dashed lines, correspond to the *seg_size* value that allows the provider’s segments to be contained in a single data segment.

Scalability. In this subsection, we examine the scalability of the Secure Miner, focusing on its capacity to efficiently manage an increasing workload in the presence of limited memory resources. We implemented three distinct test configurations gauging runtime memory usage as variations of our motivating scenario log. In particular, we considered (I) the maximum number of events per case, (II) the number of cases $|\widehat{CID}|$, and (III) the number of provisioning organizations $|\widehat{O}|$

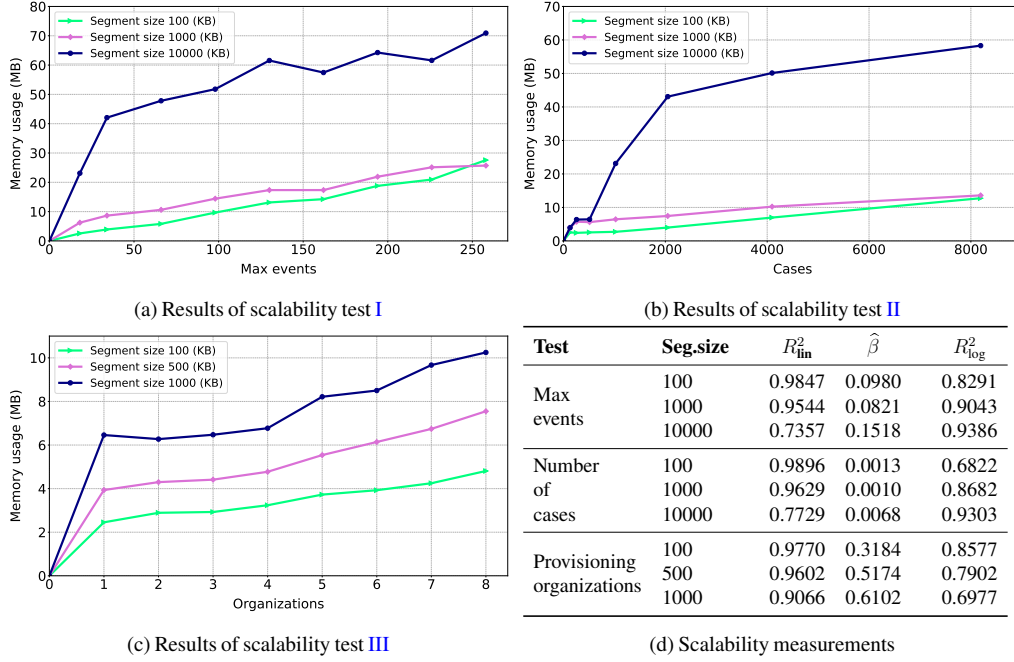


Figure 9: Scalability test results

as independent integer variables. To conduct the test on the maximum number of events, we added a loop back from the final to the initial activity of the process model, progressively increasing the number of iterations $2 \leq x_{\mathcal{U}} \leq 16$ at a step of 2, resulting in $18 + 16 \cdot (x_{\mathcal{U}} - 1)$ events. Concerning the test on the number of cases, we simulated additional process instances so that $|\widehat{CID}| = 2^{x_{cid}}$ having $x_{cid} \in \{7, 8, \dots, 13\}$. Finally, for the assessment of the number of organizations, the test necessitated the distribution of the process model activities' into a variable number of pools, each representing a different organization ($|\hat{\mathcal{O}}| \in \{1, 2, \dots, 8\}$). We parameterized the above configurations with three segment sizes (in KiloBytes): $seg_size \in \{100, 1000, 10000\}$ for tests I and II, and $seg_size \in \{100, 500, 1000\}$ for test III (the range is reduced without loss of generality to compensate the partitioning of activities into multiple organizations). To facilitate a more rigorous interpretation of the output trends across

421 varying *seg_size* configurations, we employ two well-known statistical measures.
 422 As a primary measure of goodness-of-fit, we employ the coefficient of determination
 423 R^2 [28], which assesses the degree to which the observed data adheres to the linear
 424 (R_{lin}^2) and logarithmic (R_{log}^2) regressions derived from curve fitting approximations.
 425 To further delve into the analysis of trends with a high R_{lin}^2 , we consider the slope $\hat{\beta}$
 426 of the approximated linear regression [29].

427 Table 9(d) lists the measurements we obtained. We describe them to elucidate the
 428 observed patterns. Figure 9(a) depicts the results of test I, focusing on the increase of
 429 memory utilization when the number of events in the event logs grows. We observe
 430 that the memory usage for *seg_size* 100 and 1000 (depicted by green and lilac lines,
 431 respectively) are quite similar, whereas the setting with *seg_size* 10,000 (blue line)
 432 exhibits significantly higher memory usage. For the settings with *seg_size* 100 and
 433 1000, R_{lin}^2 approaches 1, signifying an almost perfect approximation of the linear
 434 relation, against lower R_{log}^2 values. In these test settings, $\hat{\beta}$ is very low yet higher than
 435 0, thus indicating that memory usage is likely to continue increasing as the number
 436 of max events grows. The configuration with *seg_size* 10,000 yields a higher R_{log}^2
 437 value, thus suggesting a logarithmic trend, hence a greater likelihood of stabilizing
 438 memory usage growth rate as the number of maximum events increases. In Fig. 9(b),
 439 we present the results of test II, assessing the impact of the number of cases on
 440 the memory consumption. As expected, the configurations with *seg_size* set to
 441 100 and 1000 exhibit a trend of lower memory usage than settings with *seg_size*
 442 10,000. The R_{lin}^2 score of the trends with *seg_size* 100 and 1000 indicate a strong
 443 linear relationship between the dependent and independent variables compared to
 444 the trend with *seg_size* 10,000, which is better described by a logarithmic regression
 445 ($R_{log}^2 = 0.9303$). For the latter, the R_{log}^2 value is higher than the corresponding R_{lin}^2

thus suggesting that the logarithmic approximation is better suited to describe the trend. Differently from test I, the $\hat{\beta}$ score associated with the linear approximations of the trends with *seg_size* 100 and 1000 approaches 0, indicating that the growth rate of memory usage as the number of cases increases is negligible. In Fig. 9(c), we present the results of test III, on the relation between the number of organizations and the memory usage. The chart shows that memory usage trends increase as provisioning organizations increase for all three segment sizes. The R_{lin}^2 values for the three *seg_sizes* are very high, indicating a strong positive linear correlation. The test with *seg_size* 100 exhibits the slowest growth rate, as corroborated by the lowest $\hat{\beta}$ result (0.3184). For the configuration with *seg_size* 500, the memory usage increases slightly faster ($\hat{\beta} = 0.5174$). With *seg_size* 1000, the overall memory usage increases significantly faster than the previous configurations ($\hat{\beta} = 0.6102$). We derive from these findings that the Secure Miner may encounter scalability issues when handling settings with a large number of provisioning organizations. Further investigation is warranted to determine the precise cause of this behavior and identify potential mitigation strategies.

In the next section, we conclude our work and outline future research directions based upon our current findings and the limitations of our approach.

8. Conclusion and Future Work

Confidentiality is paramount in inter-organizational process mining due to the transmission of sensitive data across organizational boundaries. Our research investigates a decentralized secrecy-preserving approach that enables organizations to employ process mining techniques with event logs from multiple organizations while ensuring the protection of privacy and confidentiality. Our solution offers

470 a number of directions to walk along for improvement. We operate under the
471 assumption of fair conduct by data provisioners and do not account for the presence
472 of injected or maliciously manipulated event logs. In addition, we assume that miners
473 and provisioners exchange messages in reliable communication channels where no
474 loss or bit corruption occurs. Our approach relies on certain assumptions about event
475 log data, including the existence of a universal clock for event timestamps, which may
476 not be realistic in situations where organizations are not perfectly synchronized. We
477 aim at enhancing our approach to make it robust to the relaxation of these constraints.
478 Our future work encompasses the integration of usage control policies that specify
479 rules on event logs' utilization. We plan to design policy enforcement and monitoring
480 mechanisms to achieve this goal following the principles already addressed in [30, 31].
481 Our solution embraces process mining techniques in a general way. However, we
482 believe the presented approach is compatible with declarative model representations
483 [32]. Therefore, trusted applications could compute and store the entire set of
484 rules representing a business process, and users may interact with them via trusted
485 queries. Finally, in our implementation, we have focused on process discovery tasks.
486 Nevertheless, our approach has the potential to seamlessly cover a wider array of
487 process mining functionalities such as *conformance checking*, and *performance*
488 *analysis* techniques. Implementing them and showing their integrability with our
489 approach paves the path for future research endeavors.

490 **Acknowledgments.** The authors thank Giuseppe Ateniese for the fruitful discussion
491 and insights. This research work was partly funded by MUR under PRIN grant
492 B87G22000450001 (PINPOINT), by the Latium Region under PO FSE+ grant
493 B83C22004050009 (PPMPP), and by the EU-NGEU under the NRRP MUR grant
494 PE00000014 (SERICS).

EXTENSION PLAN:

Extended introduction

Add Background Section

Add more related work

Modify the motivating scenario (and the design alongside the implementation) with more attributes involved in the event log (for example, the famous ID that might not be shared among providers, many of those having differing attributes and codes to refer to an instance).

Add Notation and formalization of event logs, merging, partitioning and segmentation

Add Soundness and completeness theorems

Add threat model

Full pseudocode of the protocol ✓

More real-world event logs (plus two, at least) with associated tests

Add communication overhead v. segment size charts: elab time vs segment size; total time (incl. network) vs segment size.

Integrate declarative conformance checking (let it be with Janus or MINERful).

495

References

496

497 [1] W. M. P. van der Aalst, et al., Process mining manifesto, in: BPM Workshops,
498 2012, pp. 169–194.

499 [2] W. M. P. van der Aalst, Intra-and inter-organizational process mining: Dis-
500 covering processes within and between organizations, in: PoEM, 2011, pp.
501 1–11.

502 [3] C. Liu, Q. Li, X. Zhao, Challenges and opportunities in collaborative business
503 process management: Overview of recent advances and introduction to the
504 special issue, Inf. Syst. Front. 11 (2009) 201–209.

505 [4] M. Müller, N. Ostern, Koljada, et al., Trust mining: analyzing trust in
506 collaborative business processes, IEEE Access (2021) 65044–65065.

507 [5] M. Sabt, M. Achemlal, A. Bouabdallah, Trusted execution environment: What

- 508 it is, and what it is not, in: 2015 IEEE TrustCom/BigDataSE/ISPA, 2015, pp.
509 57–64.
- 510 [6] M. Müller, A. Simonet-Boulogne, S. Sengupta, O. Beige, Process mining in
511 trusted execution environments: Towards hardware guarantees for trust-aware
512 inter-organizational process analysis, in: ICPM, 2021, pp. 369–381.
- 513 [7] G. Elkoumy, S. A. Fahrenkrog-Petersen, et al., Shareprom: A tool for privacy-
514 preserving inter-organizational process mining, in: BPM (PhD/Demos), 2020,
515 pp. 72–76.
- 516 [8] G. Elkoumy, S. A. Fahrenkrog-Petersen, et al., Secure multi-party computation
517 for inter-organizational process mining, in: BPMDS/EMMSAD, 2020, pp.
518 166–181.
- 519 [9] W. M. P. van der Aalst, Federated process mining: Exploiting event data across
520 organizational boundaries, in: SMDS 2021, 2021, pp. 1–7.
- 521 [10] R. Cramer, I. Damgård, J. B. Nielsen, Secure Multiparty Computation and
522 Secret Sharing, Cambridge University Press, 2015.
- 523 [11] C. Zhao, S. Zhao, Zhao, et al., Secure multi-party computation: Theory,
524 practice and applications, Inf. Sci. 476 (2019) 357–372.
- 525 [12] J. Claes, G. Poels, Merging event logs for process mining: A rule based
526 merging method and rule suggestion algorithm, Expert Syst. Appl. 41 (2014)
527 7291–7306.
- 528 [13] J. D. Hernandez-Resendiz, E. Tello-Leal, H. M. Marin-Castro, et al., Merging
529 event logs for inter-organizational process mining, in: New Perspectives

- 530 on Enterprise Decision-Making Applying Artificial Intelligence Techniques,
531 Springer, 2021, pp. 3–26.
- 532 [14] M. Jans, M. Hosseinpour, How active learning and process mining can act as
533 continuous auditing catalyst, *Int. J. Accounting Inf. Systems* 32 (2019) 44–58.
- 534 [15] N. Koch, A. Kraus, The expressive power of UML-based web engineering, in:
535 IWWOST02, volume 16, 2002, pp. 40–41.
- 536 [16] M. Dumas, M. La Rosa, J. Mendling, H. A. Reijers, *Fundamentals of Business*
537 *Process Management, Second Edition*, Springer, 2018.
- 538 [17] V. Costan, S. Devadas, Intel SGX explained, *Cryptology ePrint Archive*
539 (2016).
- 540 [18] P. Jauernig, A.-R. Sadeghi, E. Stapf, Trusted execution environments: Proper-
541 ties, applications, and challenges, *IEEE Secur. Priv.* 18 (2020) 56–60.
- 542 [19] R. L. Rivest, A. Shamir, L. M. Adleman, A method for obtaining digital
543 signatures and public-key cryptosystems (reprint), *Commun. ACM* 26 (1983)
544 96–99.
- 545 [20] C. Cachin, R. Guerraoui, L. E. T. Rodrigues, *Introduction to Reliable and*
546 *Secure Distributed Programming (2. ed.)*, Springer, 2011.
- 547 [21] A. Kamil, G. Lowe, Understanding abstractions of secure channels, in: *FAST*,
548 2010, pp. 50–64.
- 549 [22] H. Birkholz, D. Thaler, M. Richardson, et al., Remote ATtestation procedureS
550 (RATS) Architecture, 2023.

- 551 [23] S. A. Thomas, *SSL and TLS Essentials: Securing the Web*, Wiley, 2000.
- 552 [24] A. J. M. M. Weijters, W. M. P. van der Aalst, A. K. Alves De Medeiros, *Process*
553 *mining with the HeuristicsMiner algorithm*, 2006.
- 554 [25] F. Mannhardt, *Sepsis cases - event log*, 2016. doi:[10.4121/UUID:](https://doi.org/10.4121/UUID:915D2BFB-7E84-49AD-A286-DC35F063A460)
555 [915D2BFB-7E84-49AD-A286-DC35F063A460](https://doi.org/10.4121/UUID:915D2BFB-7E84-49AD-A286-DC35F063A460).
- 556 [26] W. Steeman, *BPI challenge 2013, incidents*, 2013. doi:[10.4121/UUID:](https://doi.org/10.4121/UUID:500573E6-ACCC-4B0C-9576-AA5468B10CEE)
557 [500573E6-ACCC-4B0C-9576-AA5468B10CEE](https://doi.org/10.4121/UUID:500573E6-ACCC-4B0C-9576-AA5468B10CEE).
- 558 [27] W. M. P. van der Aalst, *Verification of workflow nets*, in: *ICATPN*, 1997, pp.
559 407–426.
- 560 [28] J. P. Barrett, *The coefficient of determination—some limitations*, *The American*
561 *Statistician* (1974) 19–20.
- 562 [29] N. Altman, M. Krzywinski, *Simple linear regression*, *Nature Methods* (2015)
563 999–1000.
- 564 [30] D. Basile, C. Di Ciccio, V. Goretti, S. Kirrane, *Blockchain based resource*
565 *governance for decentralized web environments*, *Frontiers in Blockchain*
566 (2023) 1141909.
- 567 [31] D. Basile, C. Di Ciccio, V. Goretti, S. Kirrane, *A blockchain-driven architecture*
568 *for usage control in solid*, in: *ICDCSW*, 2023, pp. 19–24.
- 569 [32] C. Di Ciccio, M. Montali, *Declarative process specifications: Reasoning,*
570 *discovery, monitoring*, in: *Process Mining Handbook*, Springer, 2022, pp.
571 108–152.