

Preserving Data Secrecy in Decentralized Inter-organizational Process Mining

Valerio Goretti^a, Davide Basile^a, Luca Barbaro^a, Claudio Di Ciccio^b

^a*Sapienza University of Rome, Italy,*

^b*Utrecht University, Netherlands,*

Abstract

Inter-organizational business processes involve multiple independent organizations collaborating to achieve mutual interests. Process mining techniques have the potential to allow these organizations to enhance operational efficiency, improve performance, and deepen the understanding of their business based on the recorded process event data. However, inter-organizational process mining faces substantial challenges, including topical secrecy concerns: The involved organizations may not be willing to expose their own data to run mining algorithms jointly with their counterparts or third parties. In this paper, we introduce CONFINE, a novel approach that unlocks process mining on multiple actors' process event data while safeguarding the secrecy and integrity of the original records in an inter-organizational business setting. To ensure that the phases of the presented interaction protocol are secure and that the processed information is hidden from involved and external actors alike, our approach resorts to a decentralized architecture comprised of trusted applications running in Trusted Execution Environments (TEEs). We show the feasibility of our solution by showcasing its application to a healthcare scenario and evaluating our implementation in terms of memory usage and scalability on real-world event logs.

Keywords: Collaborative information

1. Introduction

In today’s business landscape, organizations constantly seek ways to enhance operational efficiency, increase performance, and gain valuable insights to improve their processes. Process mining offers techniques to discover, monitor, and improve business processes by extracting knowledge from chronological records known as *event logs* [1]. Process-aware information systems record events referring to activities and interactions within a business process. The vast majority of process mining contributions consider *intra-organizational* settings, in which business processes are executed inside individual organizations. However, organizations increasingly recognize the value of collaboration and synergy in achieving operational excellence. *Inter-organizational* business processes involve several independent organizations cooperating to achieve a shared objective [2]. Despite the advantages of transparency, performance optimization, and benchmarking that companies can gain, inter-organizational process mining raises challenges that hinder its application. The major issue concerns confidentiality. Companies are reluctant to share private information required to execute process mining algorithms with their partners [3]. Indeed, letting sensitive operational data traverse organizational boundaries introduces concerns about data privacy, security, and compliance with internal regulations [4]. *Trusted Execution Environments* (TEEs) [5] can serve as fundamental enablers to balance the need for insights with the need to protect sensitive information in inter-organizational settings. TEEs offer secure contexts that guarantee code integrity and data confidentiality before, during, and after its utilization.

In this paper, we propose CONFINE, a novel approach and tool aimed at enhanc-

ing collaborative information system architectures with secrecy-preserving process mining capabilities in a decentralized fashion. It resorts to *trusted applications* running in TEEs to preserve the secrecy and integrity of shared data. To pursue this aim, we design a decentralized architecture for a four-staged protocol: (i) The initial exchange of preliminary metadata, (ii) the attestation of the miner entity, (iii) the secure transmission and privacy-preserving merge of encrypted information segments amid multiple parties, (iv) the isolated and verifiable computation of process discovery algorithms on joined data. We evaluate our proof-of-concept implementation against synthetic and real-world-based data with a convergence test followed by experiments to assess the scalability of our approach.

The remainder of this paper is as follows. [Sect. 2](#) provides an overview of related work. In [Sect. 3](#), we introduce a motivating use-case scenario in healthcare. We present the CONFINE approach in [Sect. 4](#). We describe the implementation of our approach in [Sect. 5](#). In [Sect. 6](#), we report on the efficacy and efficiency tests for our solution. Finally, we conclude our work and outline future research directions in [Sect. 7](#).

2. Related Work

Despite the relative recency of this research branch across process mining and collaborative information systems, scientific literature already includes noticeable contributions to inter-organizational process mining. The work of Müller et al. [6] focuses on data privacy and security within third-party systems that mine data generated from external providers on demand. To safeguard the integrity of data earmarked for mining purposes, their research introduces a conceptual architecture that entails the execution of process mining algorithms within a cloud service environment, fortified with Trusted Execution Environments. Drawing inspiration from this foun-

48 dational contribution, our research work seeks to design a decentralized approach
49 characterized by organizational autonomy in the execution of process mining algo-
50 rithms, devoid of synchronization mechanisms involvement taking place between the
51 involved parties. A notable departure from the framework of Müller et al. lies in the
52 fact that here each participating organization retains the discretion to choose when
53 and how mining operations are conducted. Moreover, we bypass the idea of fixed
54 roles, engineering a peer-to-peer scenario in which organizations can simultaneously
55 be data provisioners or miners. Elkoumy et al. [7, 8] present Shareprom. Like our
56 work, their solution offers a means for independent entities to execute process mining
57 algorithms in inter-organizational settings while safeguarding their proprietary input
58 data from exposure to external parties operating within the same context. Share-
59 prom’s functionality, though, is confined to the execution of operations involving
60 event log abstractions [9] represented as directed acyclic graphs, which the parties
61 employ as intermediate pre-elaboration to be fed into secure multiparty computation
62 (SMPC) [10]. As the authors remark, relying on this specific graph representation
63 imposes constraints that may prove limiting in various process mining scenarios. In
64 contrast, our approach allows for the secure, ciphered transmission of event logs to
65 process mining nodes as a whole. Moreover, SMPC-based solutions require compu-
66 tationally intensive operations and synchronous cooperation among multiple parties,
67 which make these protocols challenging to manage as the number of participants
68 scales up [11]. In our research work, individual computing nodes run the calculations,
69 thus not requiring synchronization with other machines once the input data is loaded.

70 We are confronted with the imperative task of integrating event logs originat-
71 ing from different data sources and reconstructing consistent traces that describe
72 collaborative process executions. Consequently, we engage in an examination of

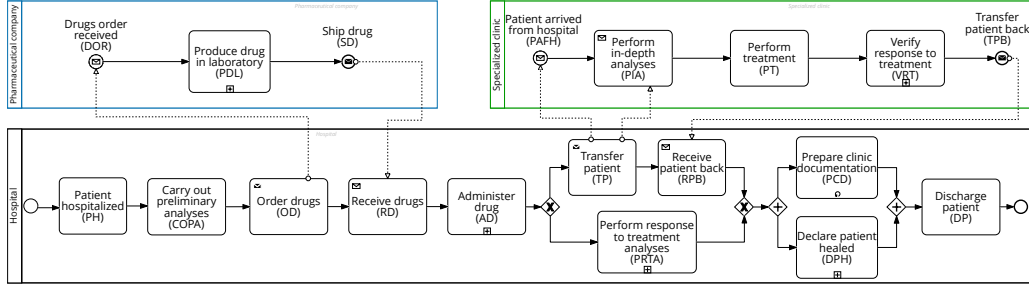


Figure 1: A BPMN collaboration diagram of a simplified healthcare scenario

methodologies delineated within the literature, each of which offers insights into the merging of event logs within inter-organizational settings. The work of Claes et al. [12] holds particular significance for our research efforts. Their seminal study introduces a two-step mechanism operating at the structured data level, contingent upon the configuration and subsequent application of merging rules. Each such rule indicates the relations between attributes of the traces and/or the activities that must hold across distinct traces to be combined. In accordance with their principles, our research incorporates a structured data-level merge based on case references and timestamps as merging attributes. The research by Hernandez et al. [13] posits a methodology functioning at the raw data level. Their approach represents traces and activities as *bag-of-words* vectors, subject to cosine similarity measurements to discern links and relationships between the traces earmarked for combination. An appealing aspect of this approach lies in its capacity to generalize the challenge of merging without necessitating a-priori knowledge of the underlying semantics inherent to the logs under consideration. However, it entails computational overhead in the treatment of data that can interfere with the overall effectiveness of our approach.

Table 1: Events from cases 312 (Alice) and 711 (Bob) recorded by the hospital, the specialized clinic, and the pharmaceutical company

Hospital						Pharmaceutical company			Specialized clinic		
Case	Timestamp	Activity	Case	Timestamp	Activity	Case	Timestamp	Activity	Case	Timestamp	Activity
312	2022-07-14T10:36	PH	312	2022-07-15T22:06	TP	312	2022-07-15T09:06	DOR	312	2022-07-16T00:06	PAFH
312	2022-07-14T16:36	COPA	711	2022-07-16T00:55	PRTA	711	2022-07-15T09:30	DOR	312	2022-07-16T01:06	PIA
711	2022-07-14T17:21	PH	711	2022-07-16T00:55	PCD	312	2022-07-15T11:06	PDL	312	2022-07-16T03:06	PT
312	2022-07-14T17:36	OD	711	2022-07-16T02:55	DPH	711	2022-07-15T11:30	PDL	312	2022-07-16T04:06	VRT
711	2022-07-14T23:21	COPA	711	2022-07-16T04:55	DP	312	2022-07-15T13:06	SD	312	2022-07-16T05:06	TPB
711	2022-07-15T00:21	OD	312	2022-07-16T07:06	RPB	711	2022-07-15T13:30	SD			
711	2022-07-15T18:55	RD	312	2022-07-16T09:06	DPH						
312	2022-07-15T19:06	RD	312	2022-07-16T09:06	PCD						
711	2022-07-15T20:55	AD	312	2022-07-16T11:06	DP						
312	2022-07-15T21:06	AD									

$$T_{312} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, PAFH, PIA, PT, VRT, TPB, RPB, DPH, PCD, DP} \rangle$$

$$T_{711} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, DPH, PCD, DP} \rangle$$

3. Motivating Scenario

For our motivating scenario, we focus on a simplified hospitalization process for the treatment of rare diseases. The process model is depicted as a BPMN diagram in Fig. 1 and involves the cooperation of three parties: a hospital, a pharmaceutical company, and a specialized clinic. For the sake of simplicity, we describe the process through two cases, recorded by the information systems as in Table 2. Alice's journey (case 312) begins when she enters the hospital for the preliminary examinations (patient hospitalized, PH). The hospital then places an order for the drugs (OD) to the pharmaceutical company for treating Alice's specific condition. Afterwards, the pharmaceutical company acknowledges that the drugs order is received (DOR), proceeds to produce the drugs in the laboratory (PDL), and ships the drugs (SD) back to the hospital. Upon receiving the medications, the hospital administers the drug (AD), and conducts an assessment to determine if Alice can be treated internally. If specialized care is required, the hospital transfers the patient (TP) to the specialized clinic. When the patient arrives from the hospital (PAFH), the specialized clinic performs in-depth analyses (PIA) and proceeds with the treatment (PT). Once the specialized clinic had completed the evaluations and verified the response to the treatment (VRT), it transfers the patient back (TPB). The hospital receives the patient

back (RPB) and prepares the clinical documentation (PCD). If Alice has successfully recovered, the hospital declares the patient as healed (DPH). When Alice's treatment is complete, the hospital discharges the patient (DP). Bob (**case 711**) enters the hospital a few hours later than Alice. His hospitalization process is similar to Alice's. However, he does not need specialized care, and his case is only treated by the hospital. Therefore, the hospital performs the response to treatment analyses (PRTA) instead of transferring him to the specialized clinic.

Both the National Institute of Statistics of the country in which the three organizations reside and the University that hosts the hospital wish to uncover information on this inter-organizational process for reporting and auditing purposes [14] via process analytics. The involved organizations share the urge for such an analysis and wish to be able to repeat the mining task also in-house. The hospital, the specialized clinic, and the pharmaceutical company have a partial view of the overall unfolding of the inter-organizational process as they record the events stemming from the parts of their pertinence. In Table 1, we show the cases 312 and 711 and the corresponding traces recorded by the hospital (i.e., T_{312}^H and T_{711}^H), the specialized clinic (i.e., T_{312}^S and T_{711}^S), and the pharmaceutical company (i.e., T_{312}^C and T_{711}^C). Those traces are projections of the two combined ones for the whole inter-organizational process: $T_{312} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, PAFH, PIA, PT, VRT, TPB, RPB, DPH, PCD, DP} \rangle$ and $T_{711} = \langle \text{PH, COPA, OD, DOR, PDL, SD, RD, AD, TP, DPH, PCD, DP} \rangle$. Results stemming from the analysis of the local cases would not provide a full picture. Data should be merged. However, to preserve the privacy of the people involved and safeguard the confidentiality of the information, the involved parties cannot give open access to their traces to other organizations. The diverging interests (being able to conduct process mining on data from multiple sources without

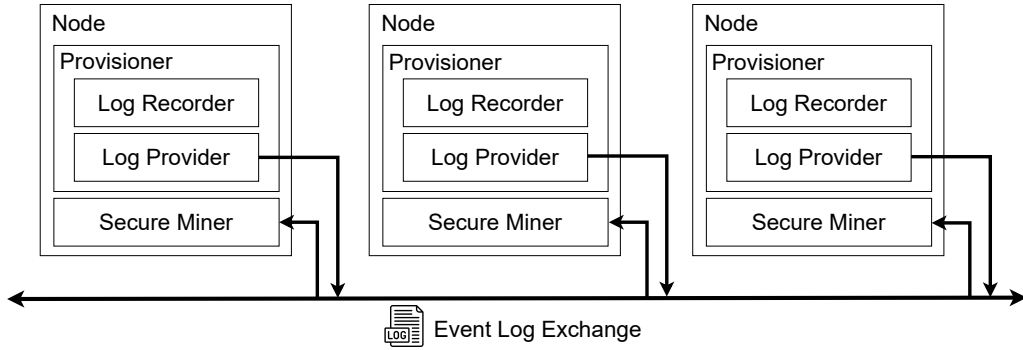


Figure 2: The CONFINE high-level architecture

132 giving away the local event logs in-clear) motivate our research. In the following,
 133 we describe the design of our solution.

134 4. Design

135 In this section, we present the high-level architecture of the CONFINE framework.
 136 We consider the main functionalities of each component, avoiding details on the
 137 employed technologies discussed in the next sections.

138 **The CONFINE architecture at large.** Our ar-
 139 chitecture involves different information systems
 140 running on multiple machines. An organiza-
 141 tion can take at least one of the following roles:
 142 **provisioning** if it delivers local event logs to
 143 be collaboratively mined; **mining** if it applies




Figure 3: Sub-components of the Secure Miner

144 process mining algorithms using event logs retrieved from provisioners. In Fig. 2,
 145 we propose the high-level schematization of the CONFINE framework. In our
 146 solution, every organization hosts one or more nodes hosting components (the
 147 names of which will henceforth be formatted with a teletype font). Depending
 148 on the played role, nodes come endowed with a Provisioner or a Secure Miner

149 component, or both. The Provisioner component consists of the following two
150 main sub-components. The Log Recorder registers the events taking place in the
151 organizations' systems. The Log Provider delivers on-demand data to mining
152 players. The hospital and all other parties in our example record Alice and Bob's
153 cases using the Log Recorder. The Log Recorder is queried by the Log Provider
154 for event logs to be made available for mining. The latter controls access to local
155 event logs by authenticating data requests by miners and rejecting those that come
156 from unauthorized parties. In our motivating scenario, the specialized clinic, the
157 pharmaceutical company, and the hospital leverage Log Providers to authenticate
158 the miner party before sending their logs. The Secure Miner component shelters
159 external event logs inside a protected environment to preserve data confidentiality
160 and integrity. Notice that Log Providers accept requests issued solely by Secure
161 Miners. Next, we provide an in-depth focus on the latter.

162 **The Secure Miner.** The primary objective of the Secure Miner is to allow
163 miners to securely execute process mining algorithms using event logs retrieved
164 from provisioners such as the specialized clinic, pharmaceutical company, and the
165 hospital in our example. Secure Miners are isolated components that guarantee
166 data inalterability and confidentiality. In [Fig. 3](#), we show a schematization of
167 the Secure Miner, which consists of four sub-components: (i) Log Requester;
168 (ii) Log Receiver; (iii) Log Manager; (iv) Log Elaborator. The Log Requester
169 and the Log Receiver are the sub-components that we employ during the event log
170 retrieval. Log Requesters send authenticable data requests to the Log Providers.
171 The Log Receiver collects event logs sent by Log Providers and entrusts them
172 to the Log Manager, securing them from accesses that are external to the Secure
173 Miner. Miners of our motivating scenario, such as the university and the national




Figure 4: UML deployment diagram of the CONFINE architecture

174 institute of statistics, employ these three components to retrieve and store Alice
 175 and Bob’s data. The Log Elaborator merges the event data locked in the Secure
 176 Miner to have a global view of the inter-organizational process comprehensive of
 177 activities executed by each involved party. Thereupon, it executes process mining
 178 algorithms in a protected environment, inaccessible from the outside computation
 179 environment. In our motivating scenario, the Log Elaborator combines the traces
 180 associated with the cases of Alice (i.e., T_{312}^H , T_{312}^S , and T_{312}^C) and Bob (i.e., T_{711}^H , T_{711}^S ,
 181 and T_{711}^C), generates the chronologically sorted traces T_{312} and T_{711} , and feeds them
 182 into the mining algorithms (see the bottom-right quadrant of Table 1).

183 5. Realization

184 In this section, we outline the technical aspects concerning the realization of
 185 our solution. Therefore, we first present the enabler technologies through which
 186 we instantiate the design principles presented in Sect. 4. After that, we discuss the
 187 CONFINE interaction protocol. Finally, we show the implementation details.

188 5.1. Deployment

189 [Figure 4](#) depicts a UML deployment diagram [15] to illustrate the employed tech-
190 nologies and computation environments. We recall that the Miner and Provisioner
191 nodes are drawn as separated, although organizations can host both. In our motivating
192 scenario, e.g., the hospital can be equipped with machines aimed for both mining
193 and provisioning.

194 Provisioner Nodes host the Provisioner’s components, encompassing the
195 Log Recorder and the Log Provider. The Process-Aware Information System
196 (PAIS) manifests the Log Recorder [16]. The PAIS grants access to the Log
197 Server, enabling it to retrieve event log data. The Log Server, on the other hand,
198 embodies the functionalities of the Log Provider, implementing web services
199 aimed at handling remote data requests and providing event log data to miners.

200 The Miner Node is characterized by two distinct *execution environments*: the
201 Operating System (OS) and the Trusted Execution Environment (TEE) [5]. TEEs
202 establish isolated contexts separate from the OS, safeguarding code and data through
203 hardware-based encryption mechanisms. This technology relies on dedicated
204 sections of a CPU capable of handling encrypted data within a reserved section
205 of the main memory [17]. By enforcing memory access restrictions, TEEs aim to
206 prevent one application from reading or altering the memory space of another, thus
207 enhancing the overall security of the system. This dedicated areas in memory are,
208 however, limited. Once the limits are exceeded, TEEs have to scout around in outer
209 memory areas, thus conceding the opportunity to malicious reader to understand
210 the saved data based on the memory reads and writes. To avoid this risk, TEE
211 implementations often raise errors that halt the program execution when the memory
212 demand goes beyond the available space. Therefore, the design of secure systems

213 that resort to TEEs must take into account that memory consumption must be kept
214 under control. We leverage the security guarantees provided by TEEs [18] to protect
215 a Trusted App responsible for fulfilling the functions of the Secure Miner and
216 its associated sub-components. The TEE ensures the integrity of the Trusted App
217 code, protecting it against potential malicious manipulations and unauthorized
218 access by programs running within the Operating System. Additionally, we utilize
219 the isolated environment of TEEs to securely store event log data (e.g., Alice and
220 Bob’s cases). The TEE retains a private key in the externally inaccessible section of
221 memory, paired with a public key in a Rivest-Shamir-Adleman (RSA) [19] scheme
222 for attestation (only the owner of the private key can sign messages that are verifiable
223 via the public key) and secure message encryption (only the owner of the private key
224 can decode messages that are encrypted with the corresponding public key). In our
225 solution, access to data located in the TEE is restricted solely to the Trusted App.
226 Users interact with the Trusted App through the Trusted App Interface, which
227 serves as the exclusive communication channel. The Trusted App offers secure
228 methods, invoked by the Trusted App Interface, for safely receiving information
229 from the Operating System and outsourcing the results of computations.

230 5.2. The *CONFINE* protocol

231 We orchestrate the interaction of the components in CONFINE via a protocol.

Assumptions. “We assume that every process executes the code triggered by events in a mutually exclusive way. This means that the same process does not handle two events concurrently. Once the handling of an event is terminated, the process keeps on checking if any other event is triggered. This periodic checking is assumed to be fair, and is achieved in an implicit way: it is not visible in the pseudo code we describe.” [20] “(Regular) reliable broadcast, (d’ora in poi solo Reliable broadcast) : Safety. Integrity (No Duplication, No Creation): per qualsiasi messaggio m , ogni processo corretto consegna m al più una volta, e solo se m è stato precedentemente inviato in broadcast da un processo mittente Liveness: Validity: se un processo corretto invia in broadcast un messaggio m , allora tutti i processi corretti alla fine consegnano m . Agreement : se un processo corretto consegna un messaggio m , allora tutti i processi corretti alla fine consegnano m .” I provisioner e il miner devono essere corretti, i.e., they do not crash. [20]

We resort to the following: “a general-purpose secure transport layer protocol, such as SSL/TLS [21]. The secure transport protocol provides a secure channel to the application layer, i.e., it provides a communication channel with some extra security services; for example, the secure channel may prevent an attacker from faking messages, hijacking messages (redirecting them to an unintended recipient, or reascribing them so that they seem to have come from a different sender), or learning the content of messages.” [22]

232

233 Our protocol involves two main actors: a Secure Miner (henceforth, \mathcal{M}) and
 234 one to many Log Provisioners ($\mathcal{P}_1, \dots, \mathcal{P}_n \in \hat{\mathcal{P}}$).

235 Given a finite set of events \hat{E} and a total-order relation \leq subset of $\hat{E} \times \hat{E}$, we
 236 identify an event log as the totally ordered set (\hat{E}, \leq) . In the example, ... Let
 237 \widehat{CID} be a finite non-empty set of symbols such that $|\widehat{CID}| \leq |\hat{E}|$. We assume that
 238 every event be associated with a *case identifier* $cid \in \widehat{CID}$ via a total surjective
 239 function $\text{id} : \hat{E} \rightarrow \widehat{CID}$ such that the restriction $<_{cid} = \leq \cap \{e \in \hat{E} : \text{id}(e) = cid\}^2$ of
 240 total order \leq on all events mapped to the same cid is strict (i.e., if $e \leq e'$ with $e \neq e'$
 241 and $\text{id}(e) = \text{id}(e')$ then $e' \not\leq e$). In the example, ... In other words, id acts as an
 242 equivalence relation partitioning \hat{E} into $\{\hat{E}_{cid}\}_{cid \in \widehat{CID}} \subseteq 2^{\hat{E}}$ based on the cid to which
 243 the events $e \in \hat{E}_{cid}$ map, and imposing that events are linearly ordered by the restriction
 244 of \leq on every \hat{E}_{cid} . Every pair $(\hat{E}_{cid}, <_{cid})$ thus represents a finite linearly totally
 245 ordered set (or *loset* for brevity) with $\hat{E}_{cid} \subseteq \hat{E}$ and $<_{cid} \subseteq \hat{E}_{cid} \times \hat{E}_{cid} \subseteq \leq \subseteq \hat{E} \times \hat{E}$. Let
 246 $(\hat{E}, <)$ be a loset and $(\hat{E}', <')$, $(\hat{E}'', <'')$ two (sub-)losets such that $\hat{E}' \cup \hat{E}'' \subseteq \hat{E}$ and
 247 $\hat{E}' \cap \hat{E}'' = \emptyset$, with $<'$ and $<''$ being the restrictions of $<$ on \hat{E}' and \hat{E}'' , respectively.

Add exam-
ple

Add exam-
ple

248 We define the order-preserving union $\oplus : \hat{E}^3 \times \hat{E}^3 \rightarrow \hat{E}^3$ of losets as follows:

249 $(\hat{E}', <') \oplus (\hat{E}'', <'') = (\hat{E}' \cup \hat{E}'', < \cap (\hat{E}' \cup \hat{E}'')^2)$. We can thus derive the notion of

250 case C_{cid} given a $cid \in \widehat{CID}$ as a loset of events mapping to the same cid and ordered

251 by the linear restriction $<$ of \leq over the events in C_{cid} : $cid = (\hat{E}_{cid}, <)$ where

252 $C_{cid} = \langle e_1, \dots, e_{|C_{cid}|} \rangle$ where $id(e_i) = cid \in \widehat{CID}$ for every i s.t. $1 \leq i \leq |C_{cid}|$ and $e_i < e_j$

253 for every $i \leq j \leq |C_{cid}|$.¹ Notice that the cardinality of \hat{C} and \widehat{CID} coincide. Events

254 are also the domain of a function $p : \hat{E} \rightarrow \hat{\mathcal{P}}$ mapping events to log provisioners. In

255 the example, . . . We shall denote with $C_{cid}^{\mathcal{P}}$ the loset consisting of every event $e \in C_{cid}$

256 such that $p(e) = \mathcal{P}$, with the restriction of the strict total order of C_{cid} on those events.

257 In the example, . . .

Gotta move this one earlier when we introduce the example with the partitioned event log. (Section 4.2??). Clarify the difference between segmentation (given a segsize, i.e., a segment of a case-part in a sublog) and partitioning (of a log into case-parts of sublogs)

258

259 . Then, prove that the pipeline of partitioning and segmentation has its inverse in

260 the union and merge for soundness. We separate it in four subsequent stages, namely

261 (i) *initialization*, (ii) *remote attestation*, (iii) *data transmission*, and (iv) *computation*.

262 These stages are depicted in Figs. 5(a), 5(b), 6(a) and 6(b), respectively. They are

263 mainly enacted by a Miner Node and n Provisioner Nodes. We assume their

264 communication channel is reliable [20] and secure [22]. In the following, we describe

265 each phase in detail.

266 **Initialization.** The objective of the initialization stage is to inform the miner about the

267 distribution of cases related to a business process among the Provisioner Nodes.

268 At the onset of this stage, the Log Requester within the Trusted App issues n

269 requests, one per Log Server component, to retrieve the list of case references they

Continue here re-vising the definition of order preserving union

Add example

Add example

¹We employ the angular-bracket notation here for the sake of simplicity, although it is typically used for sequences. Unlike sequences, cases do not allow for the same event to occur more than once.

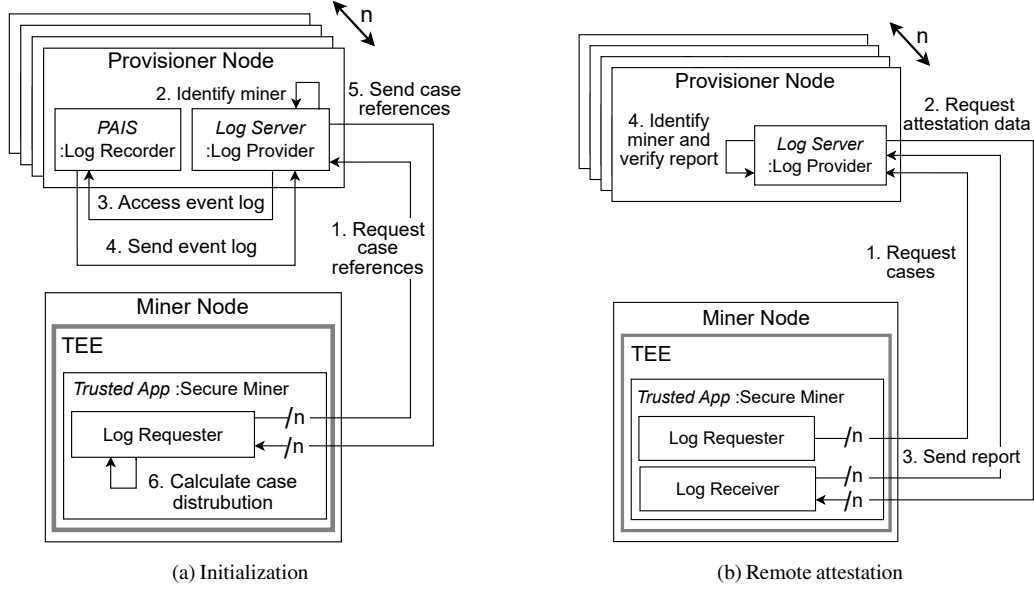


Figure 5: Unfolding example for the initialization, remote attestation phases of the CONFINE protocol

record (step 1 in Fig. 5(a)). Following sender authentication (2), each Log Server retrieves the local event log from the PAIS (3, 4) and subsequently responds to the Log Requester by providing a list of its associated case references (5). After collecting these n responses, the Log Requester delineates the distribution of cases. In the context of our motivating scenario, by the conclusion of the initialization, the miner gains knowledge that the case associated with Bob, synthesized in the traces T_{711}^H and T_{711}^C , is exclusively retained by the hospital and the specialized clinic. In contrast, the traces of Alice’s case, denoted as T_{312}^H , T_{312}^C , and T_{312}^S , are scattered across all three organizations.

Remote attestation. The remote attestation serves the purpose of establishing trust between miners and provisioners in the context of fulfilling data requests. This phase adheres to the overarching principles outlined in the RATS RFC standard [23] serving

282 as the foundation for several TEE attestation schemes (e.g., Intel EPID,² and AMD
283 SEV-SNP³). Remote attestation has a dual objective: (i) to furnish provisioners with
284 compelling evidence that the data request for an event log originates from a Trusted
285 App running within a TEE; (ii) to confirm the specific nature of the Trusted App
286 as an authentic Secure Miner software entity. This phase is triggered when the
287 Log Requester sends a new case request to the Log Server, specifying: (i) the
288 segment size (henceforth, *seg_size*), and (ii) the set of the requested case references.
289 Both parameters will be used in the subsequent *data transmission* phase. Each of
290 the n Log Servers commences the verification process by requesting the necessary
291 information from the Log Receiver to conduct the attestation (2). Subsequently, the
292 Log Receiver generates the attestation report containing the so-called *measurement*
293 of the Trusted App, which is defined as the hash value of the combination of its
294 source code and data. Once this report is signed using the attestation private key
295 associated with the TEE's hardware of the Miner Node, it is transmitted by the Log
296 Receiver to the Log Servers alongside the attestation public key of the Miner
297 Node (3). The Log Servers authenticate the miner using the public key and decrypt
298 the report (4). In this last step, the Log Servers undertake a comparison procedure
299 in which they juxtapose the measurement found within the decrypted report against
300 a predefined reference value associated with the source code of the Secure Miner.
301 If the decrypted measurement matches the predefined value, the Miner Node gains
302 trust from the provisioner.

303 **Data transmission.** Once the trusted nature of the Trusted App is verified, the Log

²sgx101.gitbook.io/sgx101/sgx-bootstrap/attestation. Accessed: 09/01/2024.

³amd.com/en/processors/amd-secure-encrypted-virtualization. Accessed: 09/01/2024.

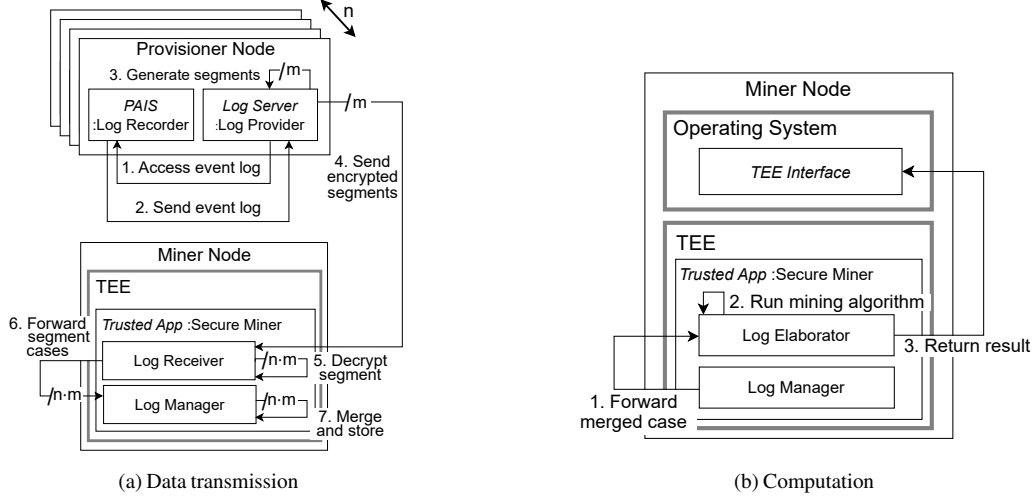


Figure 6: Unfolding example for the data transmission and computation phases of the CONFINE protocol

304 Servers proceed with the transmission of their cases. To accomplish this, each Log
 305 Server retrieves the event log from the PAIS (steps 1 and 2, in Fig. 6(a)), and filters
 306 it according to the case reference set specified by the miner. Given the constrained
 307 workload capacity of the TEE, it is imperative for Log Servers to partition the filtered
 308 event log into distinct segments. Consequently, each Log Server generates m log
 309 segments comprising a variable count of entire cases (3). The cumulative size of
 310 these segments is governed by the threshold parameter specified by the miner in the
 311 initial request (step 1 of the remote attestation phase, Fig. 5(b)). As an illustrative
 312 example from our motivating scenario, the Log Server of the hospital may structure
 313 the segmentation such that T_{312}^H and T_{711}^H reside within the same segment, whereas
 314 the specialized clinic might have T_{312}^S and T_{711}^S in separate segments. Subsequently,
 315 the n Log Servers transmit their m encrypted segments to the Log Receiver of
 316 the Trusted App (4). The Log Receiver, in turn, collects the $n \times m$ responses in a
 317 queue, processing them one at a time. After decrypting the processed segment (5),

the Log Receiver forwards the cases contained in the segment to the Log Manager (6). To reconstruct the process instance, cases belonging to the same process instance must be merged by the Log Manager resulting in a single trace (e.g., T_{312} for Alice case) comprehensive of all the events in the partial traces (e.g., T_{312}^H , T_{312}^S and T_{312}^C for Alice case). During this operation, the Log Manager applies a specific *merging schema* (i.e., a rule specifying the attributes that link two cases during the merge) as stated in [12]. In our illustrative scenario, the merging schema to combine the cases of Alice is contingent upon the linkage established through their case identifier (i.e., 312). We underline that our proposed solution facilitates the incorporation of diverse merging schemas encompassing distinct trace attributes. The outcomes arising from merging the cases within the processed segment are securely stored by the Log Manager in the TEE.

Computation. The Trusted App requires all the provisioners to have delivered cases referring to the same process instances. For example, when the hospital and the other organizations have all delivered their information concerning case 312 to the Trusted App, the process instance associated with Alice becomes eligible for computation. Upon meeting this condition, the Log Manager forwards the case earmarked for computation to the Log Elaborator (step 1 in Fig. 6(b)). Subsequently, the Log Elaborator proceeds to input the merged case into the process mining algorithm (2). Ultimately, the outcome of the computation is relayed by the Log Elaborator from the TEE to the TEE Interface running atop the Operating System of the Miner Node (3). The CONFINE protocol does not impose restrictions on the post-computational handling of results. In our motivating scenario, the University and the National Institute of Statistics, serving as miners, disseminate the outcomes of computations, generating analyses that benefit the provisioners (though the original

Algorithm 1: Secure Miner's behavior in CONFINE.

Input: $\hat{\mathcal{P}} = \{\mathcal{P}_1, \dots, \mathcal{P}_n\}$, the (references to) n log provisioners;
 seg_size , the maximum size of the log segment to be transmitted by the log provisioners;
Data: $CIDMap: \widehat{CID} \rightarrow 2^{\hat{\mathcal{P}}}$, a map from case references $cid \in \widehat{CID}$ to the set of log provisioners in $\hat{\mathcal{P}}$;
 $PMap: \hat{\mathcal{P}} \rightarrow 2^{\widehat{CID}}$, a map from log provisioners $\mathcal{P} \in \hat{\mathcal{P}}$ to the set of references to their cases in \widehat{CID} ;
 $Cases: \widehat{CID} \rightarrow \hat{\mathcal{C}}$, a map from case references $cid \in \widehat{CID}$ to a set of cases in $\hat{\mathcal{C}}$
Implements: SecureMiner \mathcal{M}
Uses: LogProvisioner \mathcal{P}

```

1  upon event  $\langle \mathcal{M}.LogRequester, INIT | \hat{\mathcal{P}}, seg\_size \rangle$  do           // The initialization phase begins here - Fig. 5(a)
2    foreach  $\mathcal{P} \in \hat{\mathcal{P}}$  do                                           // Request the list of cases provided by every Log Provider  $\mathcal{P}$ 
3      trigger  $\langle \mathcal{P}.LogProvider, CASEREFSREQ | \mathcal{M} \rangle$                //  $\mathcal{P}$  will reply with its case references (see Alg.2 line 1)
4  upon  $|\hat{\mathcal{P}}| = |\text{dom}(PMap)|$  do                                     // Once all  $\mathcal{P}$ s have answered  $\mathcal{M}$  requests their cases
5    foreach  $\mathcal{P} \in \hat{\mathcal{P}}$  do trigger  $\langle \mathcal{P}.LogProvider, CASESREQ | \mathcal{M}, seg\_size, PMap[\mathcal{P}] \rangle$  // see Alg.2 line 5

6  upon event  $\langle \mathcal{M}.LogRequester, CASEREFSRES | \mathcal{P}, CIDs \rangle$  such that  $\mathcal{P} \in \hat{\mathcal{P}}$  do //  $\mathcal{M}$  gets case references (see Alg.2 line 4)
7    foreach  $cid \in CIDs$  do                                         // For every reference to a case  $cid$ 
8       $CIDMap[cid] \leftarrow CIDMap[cid] \cup \{\mathcal{P}\}$                 // Add  $\mathcal{P}$  to the set of providers for case  $cid$  in  $CIDMap$ 
9       $PMap[\mathcal{P}] \leftarrow PMap[\mathcal{P}] \cup CIDs$                        // Register the references of the cases provided by  $\mathcal{P}$  in  $PMap$ 

10 upon event  $\langle \mathcal{M}.LogReceiver, CASESRES | \mathcal{P}, S \rangle$  such that  $\mathcal{P} \in \hat{\mathcal{P}}$  do //  $\mathcal{M}$  gets segments from  $\mathcal{P}$  (see Alg.2 line 9)
11   foreach  $C_{cid}^{\mathcal{P}} \in S$  do                                       // For every  $C_{cid}^{\mathcal{P}}$  in the delivered set of segments  $S$ , each associated with a  $cid$ 
12     if  $cid \in PMap[\mathcal{P}]$  then                                     // The store phase (see Fig. 6(a))
13        $PMap[\mathcal{P}] \leftarrow PMap[\mathcal{P}] \setminus \{cid\}$              // Remove  $cid$  from the set of cases to be provided by  $\mathcal{P}$ 
14        $CIDMap[cid] \leftarrow CIDMap[cid] \setminus \{\mathcal{P}\}$          // Remove  $\mathcal{P}$  from the set of  $cid$  providers
15        $LogManager.mergeAndStore(Cases, C_{cid}^{\mathcal{P}})$                // Update the case via  $\oplus$  and store the result in  $Cases$ 

14 upon  $CIDMap[cid] = \emptyset$  for some  $cid \in \text{dom}(CIDMap)$  do       // When all the pieces of some  $cid$  have arrived to  $\mathcal{M}$ 
15    $\text{dom}(CIDMap) \leftarrow \text{dom}(CIDMap) \setminus \{cid\}$            // Remove  $cid$  from the  $cid$  domain which still needs to be processed
16   yield  $Cases[cid]$  to LogElaborator                             // forward the  $cid$  to the Log Elaborator for mining (Fig. 6(b))

```

Algorithm 2: Provisioner's behavior in CONFINE.

Input:
Data:
Implements: Provisioner \mathcal{P}
Uses: Secure Miner \mathcal{M}

```

1  upon event  $\langle \mathcal{P}.LogProvider, CASESREFSREQ | \mathcal{M} \rangle$  do           //  $\mathcal{P}$  receives the request for case references (see Alg.1, Line 3)
2    if  $\mathcal{P}.LogProvider.authenticateMiner(\mathcal{M})$  is successful then    // Authenticate the sender miner  $\mathcal{M}$ 
3       $CIDs \leftarrow \mathcal{P}.LogRecorder.accessCaseReferences()$        // Access the case references via Log Recorder
4      trigger  $\langle \mathcal{M}.LogRequester, CASESREFSRES | \mathcal{P}, CIDs \rangle$      //  $\mathcal{P}$  sends the case references to  $\mathcal{M}$  (see Alg.1, Line 6)

5  upon event  $\langle \mathcal{P}.LogProvider, CASESREQ | \mathcal{M}, seg\_size, CIDs \rangle$  do //  $\mathcal{P}$  gets the case request (see Alg.1, Line 5)
6    if  $\mathcal{M}.LogReceiver.getAttestationReport(\mathcal{P})$  is valid then      // The attestation phase (see Fig. 5(b))
7       $\{S_1, \dots, S_m\} \leftarrow \mathcal{P}.LogProvider.segmentEventLog(\mathcal{P}.LogRecorder.accessEventLog(CIDs), seg\_size)$  // Segment the log
8      foreach  $i \in \{1, \dots, m\}$  do                                // For every split segment ...
9        trigger  $\langle \mathcal{M}.LogReceiver, CASESRES | \mathcal{P}, S_i \rangle$        // ...  $\mathcal{P}$  sends the segment to  $\mathcal{M}$  (see Alg.1, Line 10)

```

343 data are never revealed in clear). Furthermore, our protocol enables the potential for
344 provisioners to have their proprietary Secure Miner, allowing them autonomous
345 control over the computed results.

Table 2: Event logs used for our experiments

Name	Type	Activities	Cases	Max events	Min events	Avg. events	Organization \mapsto Activities
Motivating scenario	Synthetic	19	1000	18	9	14	$\mathcal{O}^P \mapsto 3, \mathcal{O}^C \mapsto 5, \mathcal{O}^H \mapsto 14$
Sepsis [26]	Real	16	1050	185	3	15	$\mathcal{O}^1 \mapsto 1, \mathcal{O}^2 \mapsto 1, \mathcal{O}^3 \mapsto 14$
BPIC2013 [27]	Real	7	1487	123	1	9	$\mathcal{O}^1 \mapsto 6, \mathcal{O}^2 \mapsto 7, \mathcal{O}^3 \mapsto 6$

5.3. Implementation

We implemented the Secure Miner component as an Intel SGX⁴ trusted application, encoded in Go through the EGo framework.⁵ We resort to a TLS [21] communication channel between miners and provisioners over the HTTP web protocol to secure the information exchange. To demonstrate the effectiveness of our framework, we re-implemented and integrated the *HeuristicsMiner* discovery algorithm [24] within the Trusted Application. Our implementation of CONFINE, including the *HeuristicsMiner* in Go, is openly accessible at the following URL: github.com/Process-in-Chains/CONFINE/.

6. Evaluation

In this section, we evaluate our approach through the testing of our tool implementation. We begin with a convergence analysis to demonstrate the correctness of the collaborative data exchange process. Subsequently, we gauge the memory usage with synthetic and real-life event logs, to observe the trend during the enactment of our protocol and assess scalability. We recall that we focus on memory utilization since the availability of space in the dedicated areas is limited as we discussed in Sect. 5.1. We discuss our experimental results in the following. For the sake of reproducibility,

⁴sgx101.gitbook.io/sgx101/. Accessed: 09/01/2024.

⁵docs.edgeless.systems/ego. Accessed: 09/01/2024.




Figure 7: *HeuristicsMiner* output in CONFINE

we make available all the testbeds and results in our public code repository (linked above).

Output convergence. To experimentally validate the correctness of our approach in the transmission and computation phases (see Sect. 5), we run a *convergence* test. To this end, we created a synthetic event log consisting of 1000 cases of 14 events on average (see Table 2) by simulating the inter-organizational process of our motivating scenario (see Fig. 1)⁶ and we partitioned it in three sub-logs (one per involved organization), an excerpt of which is listed in Table 1. We run the stand-alone *HeuristicsMiner* on the former, and processed the latter through our CONFINE toolchain. As expected, the results converge and are depicted in Fig. 7 in the form of a workflow net [25]. For clarity, we have colored activities recorded by the organizations following the scheme of Table 2 (black for the hospital, blue for the pharmaceutical company, and green for the specialized clinic).

Memory usage. Figures 8(a) and 8(b) display plots corresponding to the runtime memory utilization of our CONFINE implementation (in MegaBytes). Differently from Fig. 8(b), Fig. 8(a) excludes the computation stage by leaving the *HeuristicsMiner* inactive so as to isolate the execution from the mining-specific

⁶We generated the event log through BIMP (<https://bimp.cs.ut.ee/>). We filtered the generated log by keeping the sole events that report on the completion of activities, and removing the start and end events of the pharmaceutical company and specialized clinic’s sub-processes.

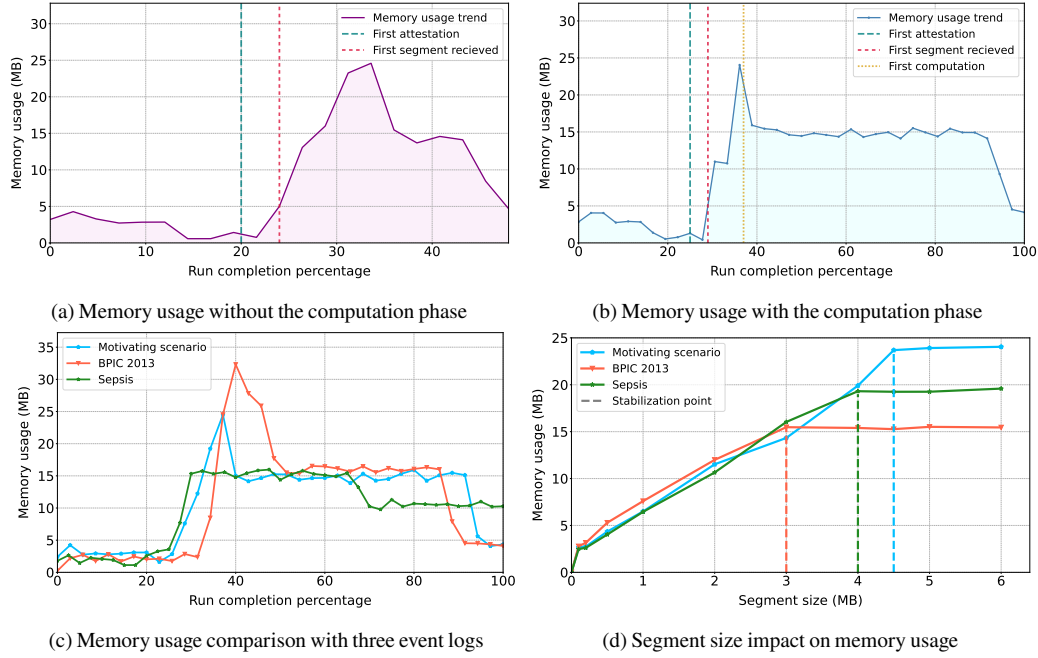


Figure 8: Memory usage test results

operations. The dashed lines mark the starting points for the remote attestation, the data transmission and the computation stages. We held the *seg_size* constant at 2000 KiloBytes. We observe that the data transmission stage reaches the highest peak of memory utilization, which is then partially freed by the subsequent computation stage, steadily occupying memory space at a lower level. To verify whether this phenomenon is due to the synthetic nature of our simulation-based event log, we also gauge the runtime memory usage of two public real-world event logs too (Sepsis [26] and BPIC 2013 [27]). The characteristics of the event logs are summarized in Table 2. Since those are *intra-organizational* event logs, we split the contents to mimic an *inter-organizational* context. In particular, we separated the Sepsis event log based on the distinction between normal-care and intensive-care paths, as if they were conducted by two distinct organizations. Similarly, we processed the BPIC 2013

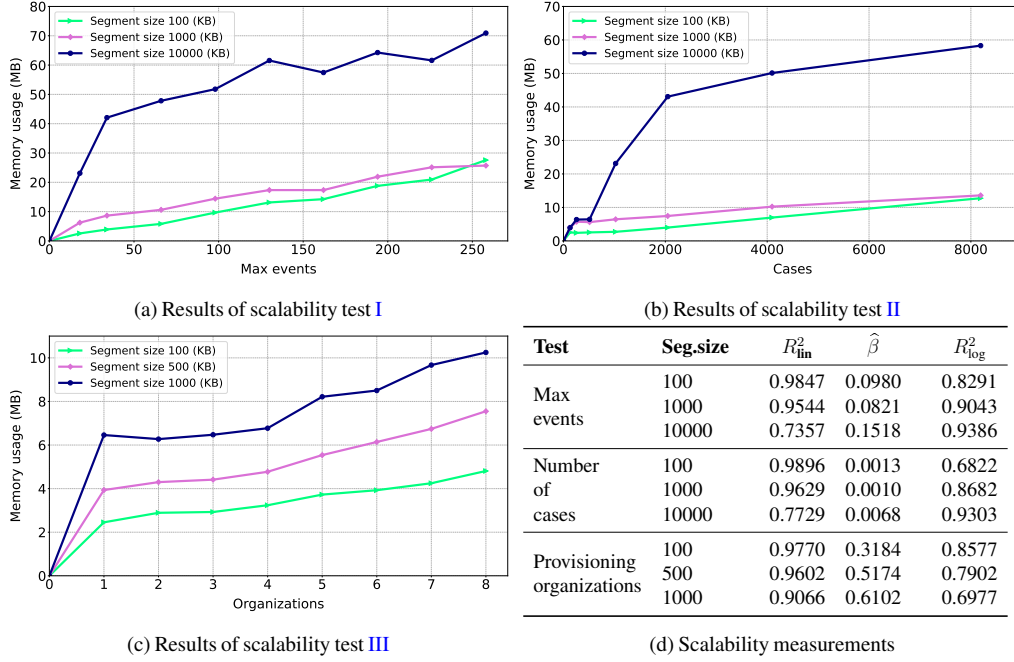


Figure 9: Scalability test results

event log to sort it out into the three departments of the Volvo IT incident management system. Figure 8(c) depicts the results. We observe that the processing of the BPIC 2013 event log demands more memory, particularly during the initial stages, probably owing to its larger size. Conversely, the Sepsis event log turns out to entail the least expensive run. To verify whether these trends are affected by the dimension of the exchanged data segments, we conducted an additional test to examine the trend of memory usage as the *seg_size* varies with all the aforementioned event logs. Notably, the polylines displayed in Fig. 8(d) indicate a linear increment of memory occupation until a breakpoint is reached. After that, the memory in use is steady. These points, marked by vertical dashed lines, correspond to the *seg_size* value that allows the provider's segments to be contained in a single data segment.

Scalability. In this subsection, we examine the scalability of the Secure Miner,

404 focusing on its capacity to efficiently manage an increasing workload in the presence
 405 of limited memory resources. We implemented three distinct test configurations
 406 gauging runtime memory usage as variations of our motivating scenario log. In
 407 particular, we considered (I) the maximum number of events per case, (II) the
 408 number of cases $|\widehat{\text{CID}}|$, and (III) the number of provisioning organizations $|\widehat{\mathcal{O}}|$
 409 as independent integer variables. To conduct the test on the maximum number of
 410 events, we added a loop back from the final to the initial activity of the process
 411 model, progressively increasing the number of iterations $2 \leq x_{\mathcal{O}} \leq 16$ at a step of 2,
 412 resulting in $18 + 16 \cdot (x_{\mathcal{O}} - 1)$ events. Concerning the test on the number of cases, we
 413 simulated additional process instances so that $|\widehat{\text{CID}}| = 2^{x_{cid}}$ having $x_{cid} \in \{7, 8, \dots, 13\}$.
 414 Finally, for the assessment of the number of organizations, the test necessitated the
 415 distribution of the process model activities' into a variable number of pools, each
 416 representing a different organization ($|\widehat{\mathcal{O}}| \in \{1, 2, \dots, 8\}$). We parameterized the above
 417 configurations with three segment sizes (in KiloBytes): $seg_size \in \{100, 1000, 10000\}$
 418 for tests I and II, and $seg_size \in \{100, 500, 1000\}$ for test III (the range is reduced
 419 without loss of generality to compensate the partitioning of activities into multiple
 420 organizations). To facilitate a more rigorous interpretation of the output trends across
 421 varying seg_size configurations, we employ two well-known statistical measures.
 422 As a primary measure of goodness-of-fit, we employ the coefficient of determination
 423 R^2 [28], which assesses the degree to which the observed data adheres to the linear
 424 (R_{lin}^2) and logarithmic (R_{log}^2) regressions derived from curve fitting approximations.
 425 To further delve into the analysis of trends with a high R_{lin}^2 , we consider the slope $\hat{\beta}$
 426 of the approximated linear regression [29].

427 Table 9(d) lists the measurements we obtained. We describe them to elucidate the
 428 observed patterns. Figure 9(a) depicts the results of test I, focusing on the increase of

429 memory utilization when the number of events in the event logs grows. We observe
 430 that the memory usage for *seg_size* 100 and 1000 (depicted by green and lilac lines,
 431 respectively) are quite similar, whereas the setting with *seg_size* 10,000 (blue line)
 432 exhibits significantly higher memory usage. For the settings with *seg_size* 100 and
 433 1000, R_{lin}^2 approaches 1, signifying an almost perfect approximation of the linear
 434 relation, against lower R_{log}^2 values. In these test settings, $\hat{\beta}$ is very low yet higher than
 435 0, thus indicating that memory usage is likely to continue increasing as the number
 436 of max events grows. The configuration with *seg_size* 10,000 yields a higher R_{log}^2
 437 value, thus suggesting a logarithmic trend, hence a greater likelihood of stabilizing
 438 memory usage growth rate as the number of maximum events increases. In Fig. 9(b),
 439 we present the results of test II, assessing the impact of the number of cases on
 440 the memory consumption. As expected, the configurations with *seg_size* set to
 441 100 and 1000 exhibit a trend of lower memory usage than settings with *seg_size*
 442 10,000. The R_{lin}^2 score of the trends with *seg_size* 100 and 1000 indicate a strong
 443 linear relationship between the dependent and independent variables compared to
 444 the trend with *seg_size* 10,000, which is better described by a logarithmic regression
 445 ($R_{log}^2 = 0.9303$). For the latter, the R_{log}^2 value is higher than the corresponding R_{lin}^2
 446 thus suggesting that the logarithmic approximation is better suited to describe the
 447 trend. Differently from test I, the $\hat{\beta}$ score associated with the linear approximations
 448 of the trends with *seg_size* 100 and 1000 approaches 0, indicating that the growth
 449 rate of memory usage as the number of cases increases is negligible. In Fig. 9(c), we
 450 present the results of test III, on the relation between the number of organizations
 451 and the memory usage. The chart shows that memory usage trends increase as
 452 provisioning organizations increase for all three segment sizes. The R_{lin}^2 values
 453 for the three *seg_sizes* are very high, indicating a strong positive linear correlation.

454 The test with *seg_size* 100 exhibits the slowest growth rate, as corroborated by the
455 lowest $\hat{\beta}$ result (0.3184). For the configuration with *seg_size* 500, the memory usage
456 increases slightly faster ($\hat{\beta} = 0.5174$). With *seg_size* 1000, the overall memory usage
457 increases significantly faster than the previous configurations ($\hat{\beta} = 0.6102$). We
458 derive from these findings that the Secure Miner may encounter scalability issues
459 when handling settings with a large number of provisioning organizations. Further
460 investigation is warranted to determine the precise cause of this behavior and identify
461 potential mitigation strategies.

462 In the next section, we conclude our work and outline future research directions
463 based upon our current findings and the limitations of our approach.

464 7. Conclusion and Future Work

465 Confidentiality is paramount in inter-organizational process mining due to the
466 transmission of sensitive data across organizational boundaries. Our research
467 investigates a decentralized secrecy-preserving approach that enables organizations
468 to employ process mining techniques with event logs from multiple organizations
469 while ensuring the protection of privacy and confidentiality. Our solution offers
470 a number of directions to walk along for improvement. We operate under the
471 assumption of fair conduct by data provisioners and do not account for the presence
472 of injected or maliciously manipulated event logs. In addition, we assume that miners
473 and provisioners exchange messages in reliable communication channels where no
474 loss or bit corruption occurs. Our approach relies on certain assumptions about event
475 log data, including the existence of a universal clock for event timestamps, which may
476 not be realistic in situations where organizations are not perfectly synchronized. We
477 aim at enhancing our approach to make it robust to the relaxation of these constraints.

478 Our future work encompasses the integration of usage control policies that specify
479 rules on event logs' utilization. We plan to design policy enforcement and monitoring
480 mechanisms to achieve this goal following the principles already addressed in [30, 31].
481 Our solution embraces process mining techniques in a general way. However, we
482 believe the presented approach is compatible with declarative model representations
483 [32]. Therefore, trusted applications could compute and store the entire set of
484 rules representing a business process, and users may interact with them via trusted
485 queries. Finally, in our implementation, we have focused on process discovery tasks.
486 Nevertheless, our approach has the potential to seamlessly cover a wider array of
487 process mining functionalities such as *conformance checking*, and *performance*
488 *analysis* techniques. Implementing them and showing their integrability with our
489 approach paves the path for future research endeavors.

490 **Acknowledgments.** The authors thank Giuseppe Ateniese for the fruitful discussion
491 and insights. This research work was partly funded by MUR under PRIN grant
492 B87G22000450001 (PINPOINT), by the Latium Region under PO FSE+ grant
493 B83C22004050009 (PPMPP), and by the EU-NGEU under the NRRP MUR grant
494 PE000000014 (SERICS).

EXTENSION PLAN:

Extended introduction

Add Background Section

Add more related work

Modify the motivating scenario (and the design alongside the implementation) with more attributes involved in the event log (for example, the famous ID that might not be shared among providers, many of those having differing attributes and codes to refer to an instance).

Add Notation and formalization of event logs, merging, partitioning and segmentation

Add Soundness and completeness theorems

Add threat model

Full pseudocode of the protocol

More real-world event logs (plus two, at least) with associated tests

Add communication overhead v. segment size charts: elab time vs segment size; total time (incl. network) vs segment size.

Integrate declarative conformance checking (let it be with Janus or MINERful).

495

References

496

497 [1] W. M. P. van der Aalst, et al., Process mining manifesto, in: BPM Workshops,
498 2012, pp. 169–194.

499 [2] W. M. P. van der Aalst, Intra-and inter-organizational process mining: Dis-
500 covering processes within and between organizations, in: PoEM, 2011, pp.
501 1–11.

502 [3] C. Liu, Q. Li, X. Zhao, Challenges and opportunities in collaborative business
503 process management: Overview of recent advances and introduction to the
504 special issue, Inf. Syst. Front. 11 (2009) 201–209.

505 [4] M. Müller, N. Ostern, Koljada, et al., Trust mining: analyzing trust in
506 collaborative business processes, IEEE Access (2021) 65044–65065.

507 [5] M. Sabt, M. Achemlal, A. Bouabdallah, Trusted execution environment: What

- 508 it is, and what it is not, in: 2015 IEEE TrustCom/BigDataSE/ISPA, 2015, pp.
509 57–64.
- 510 [6] M. Müller, A. Simonet-Boulogne, S. Sengupta, O. Beige, Process mining in
511 trusted execution environments: Towards hardware guarantees for trust-aware
512 inter-organizational process analysis, in: ICPM, 2021, pp. 369–381.
- 513 [7] G. Elkoumy, S. A. Fahrenkrog-Petersen, et al., Shareprom: A tool for privacy-
514 preserving inter-organizational process mining, in: BPM (PhD/Demos), 2020,
515 pp. 72–76.
- 516 [8] G. Elkoumy, S. A. Fahrenkrog-Petersen, et al., Secure multi-party computation
517 for inter-organizational process mining, in: BPMDS/EMMSAD, 2020, pp.
518 166–181.
- 519 [9] W. M. P. van der Aalst, Federated process mining: Exploiting event data across
520 organizational boundaries, in: SMDS 2021, 2021, pp. 1–7.
- 521 [10] R. Cramer, I. Damgård, J. B. Nielsen, Secure Multiparty Computation and
522 Secret Sharing, Cambridge University Press, 2015.
- 523 [11] C. Zhao, S. Zhao, Zhao, et al., Secure multi-party computation: Theory,
524 practice and applications, Inf. Sci. 476 (2019) 357–372.
- 525 [12] J. Claes, G. Poels, Merging event logs for process mining: A rule based
526 merging method and rule suggestion algorithm, Expert Syst. Appl. 41 (2014)
527 7291–7306.
- 528 [13] J. D. Hernandez-Resendiz, E. Tello-Leal, H. M. Marin-Castro, et al., Merging
529 event logs for inter-organizational process mining, in: New Perspectives

- 530 on Enterprise Decision-Making Applying Artificial Intelligence Techniques,
531 Springer, 2021, pp. 3–26.
- 532 [14] M. Jans, M. Hosseinpour, How active learning and process mining can act as
533 continuous auditing catalyst, *Int. J. Accounting Inf. Systems* 32 (2019) 44–58.
- 534 [15] N. Koch, A. Kraus, The expressive power of UML-based web engineering, in:
535 IWWOST02, volume 16, 2002, pp. 40–41.
- 536 [16] M. Dumas, M. La Rosa, J. Mendling, H. A. Reijers, *Fundamentals of Business*
537 *Process Management, Second Edition*, Springer, 2018.
- 538 [17] V. Costan, S. Devadas, Intel SGX explained, *Cryptology ePrint Archive*
539 (2016).
- 540 [18] P. Jauernig, A.-R. Sadeghi, E. Stapf, Trusted execution environments: Proper-
541 ties, applications, and challenges, *IEEE Secur. Priv.* 18 (2020) 56–60.
- 542 [19] R. L. Rivest, A. Shamir, L. M. Adleman, A method for obtaining digital
543 signatures and public-key cryptosystems (reprint), *Commun. ACM* 26 (1983)
544 96–99.
- 545 [20] C. Cachin, R. Guerraoui, L. E. T. Rodrigues, *Introduction to Reliable and*
546 *Secure Distributed Programming (2. ed.)*, Springer, 2011.
- 547 [21] S. A. Thomas, *SSL and TLS Essentials: Securing the Web*, Wiley, 2000.
- 548 [22] A. Kamil, G. Lowe, Understanding abstractions of secure channels, in: *FAST*,
549 2010, pp. 50–64.

- 550 [23] H. Birkholz, D. Thaler, M. Richardson, et al., Remote ATtestation procedureS
551 (RATS) Architecture, 2023.
- 552 [24] A. J. M. M. Weijters, W. M. P. van der Aalst, A. K. Alves De Medeiros, Process
553 mining with the HeuristicsMiner algorithm, 2006.
- 554 [25] W. M. P. van der Aalst, Verification of workflow nets, in: ICATPN, 1997, pp.
555 407–426.
- 556 [26] F. Mannhardt, Sepsis cases - event log, 2016. doi:[10.4121/UUID:
557 915D2BFB-7E84-49AD-A286-DC35F063A460](https://doi.org/10.4121/UUID:915D2BFB-7E84-49AD-A286-DC35F063A460).
- 558 [27] W. Steeman, BPI challenge 2013, incidents, 2013. doi:[10.4121/UUID:
559 500573E6-ACCC-4B0C-9576-AA5468B10CEE](https://doi.org/10.4121/UUID:500573E6-ACCC-4B0C-9576-AA5468B10CEE).
- 560 [28] J. P. Barrett, The coefficient of determination—some limitations, The American
561 Statistician (1974) 19–20.
- 562 [29] N. Altman, M. Krzywinski, Simple linear regression, Nature Methods (2015)
563 999–1000.
- 564 [30] D. Basile, C. Di Ciccio, V. Goretti, S. Kirrane, Blockchain based resource
565 governance for decentralized web environments, Frontiers in Blockchain
566 (2023) 1141909.
- 567 [31] D. Basile, C. Di Ciccio, V. Goretti, S. Kirrane, A blockchain-driven architecture
568 for usage control in solid, in: ICDCSW, 2023, pp. 19–24.
- 569 [32] C. Di Ciccio, M. Montali, Declarative process specifications: Reasoning,
570 discovery, monitoring, in: Process Mining Handbook, Springer, 2022, pp.
571 108–152.